

Faculty of Natural & Mathematical Sciences Careers Events



Tech Series: Infosys InStep Internship Information Session



15 Feb | 13:00 – 15:00
S-1.27, Strand Campus

Infosys®

InStep is Infosys' Global Internship Program that provides the platform to interact with a diverse group of employees and interns from across the globe. This paid internship program, seeks to provide you the opportunity to work with/explore new and emerging technologies for a duration between 8-12 weeks, while supporting you gain unparalleled international work experience.

#KCLCareers

Register for this event at King's
CareerConnect:
kcl.ac.uk/careers/events



Faculty of Natural & Mathematical Sciences Careers Events



Tech Series: GSK – Tips and experiences on getting on a fortune 500 company



1 Mar | 13:00 – 14:00

S-1.27, Strand Campus

Graduates can join us in a variety of exciting roles across several business functions. If you're enthusiastic, committed to making a difference and have a strong interest in healthcare, then we'd love to hear from you. Come and hear from two King's alumni about how they got hired at GSK and how you could too.



#KCLCareers

Register for this event at King's
CareerConnect:
kcl.ac.uk/careers/events



Faculty of Natural & Mathematical Sciences Careers Events



Tech Series: Ericsson

 **8 Mar** | 13:00 – 14:00
 **S-1.27**, Strand Campus

Over the past 140 years, Ericsson has been at the forefront of communications technology. Today, we are committed to maximizing customer value by continuously evolving our business portfolio and leading the ICT industry. Come and learn what a career with Ericsson could mean for you as well as hearing about this year's Ericsson Innovation Award. A competition open to students of all disciplines.



 [#KCLCareers](#)

Register for this event at King's
CareerConnect:
kcl.ac.uk/careers/events



Structured Query Language (SQL 2)

Queries, Update, Delete

Week 5: Database Systems (4CCS1DBS)

13 Feb 2017

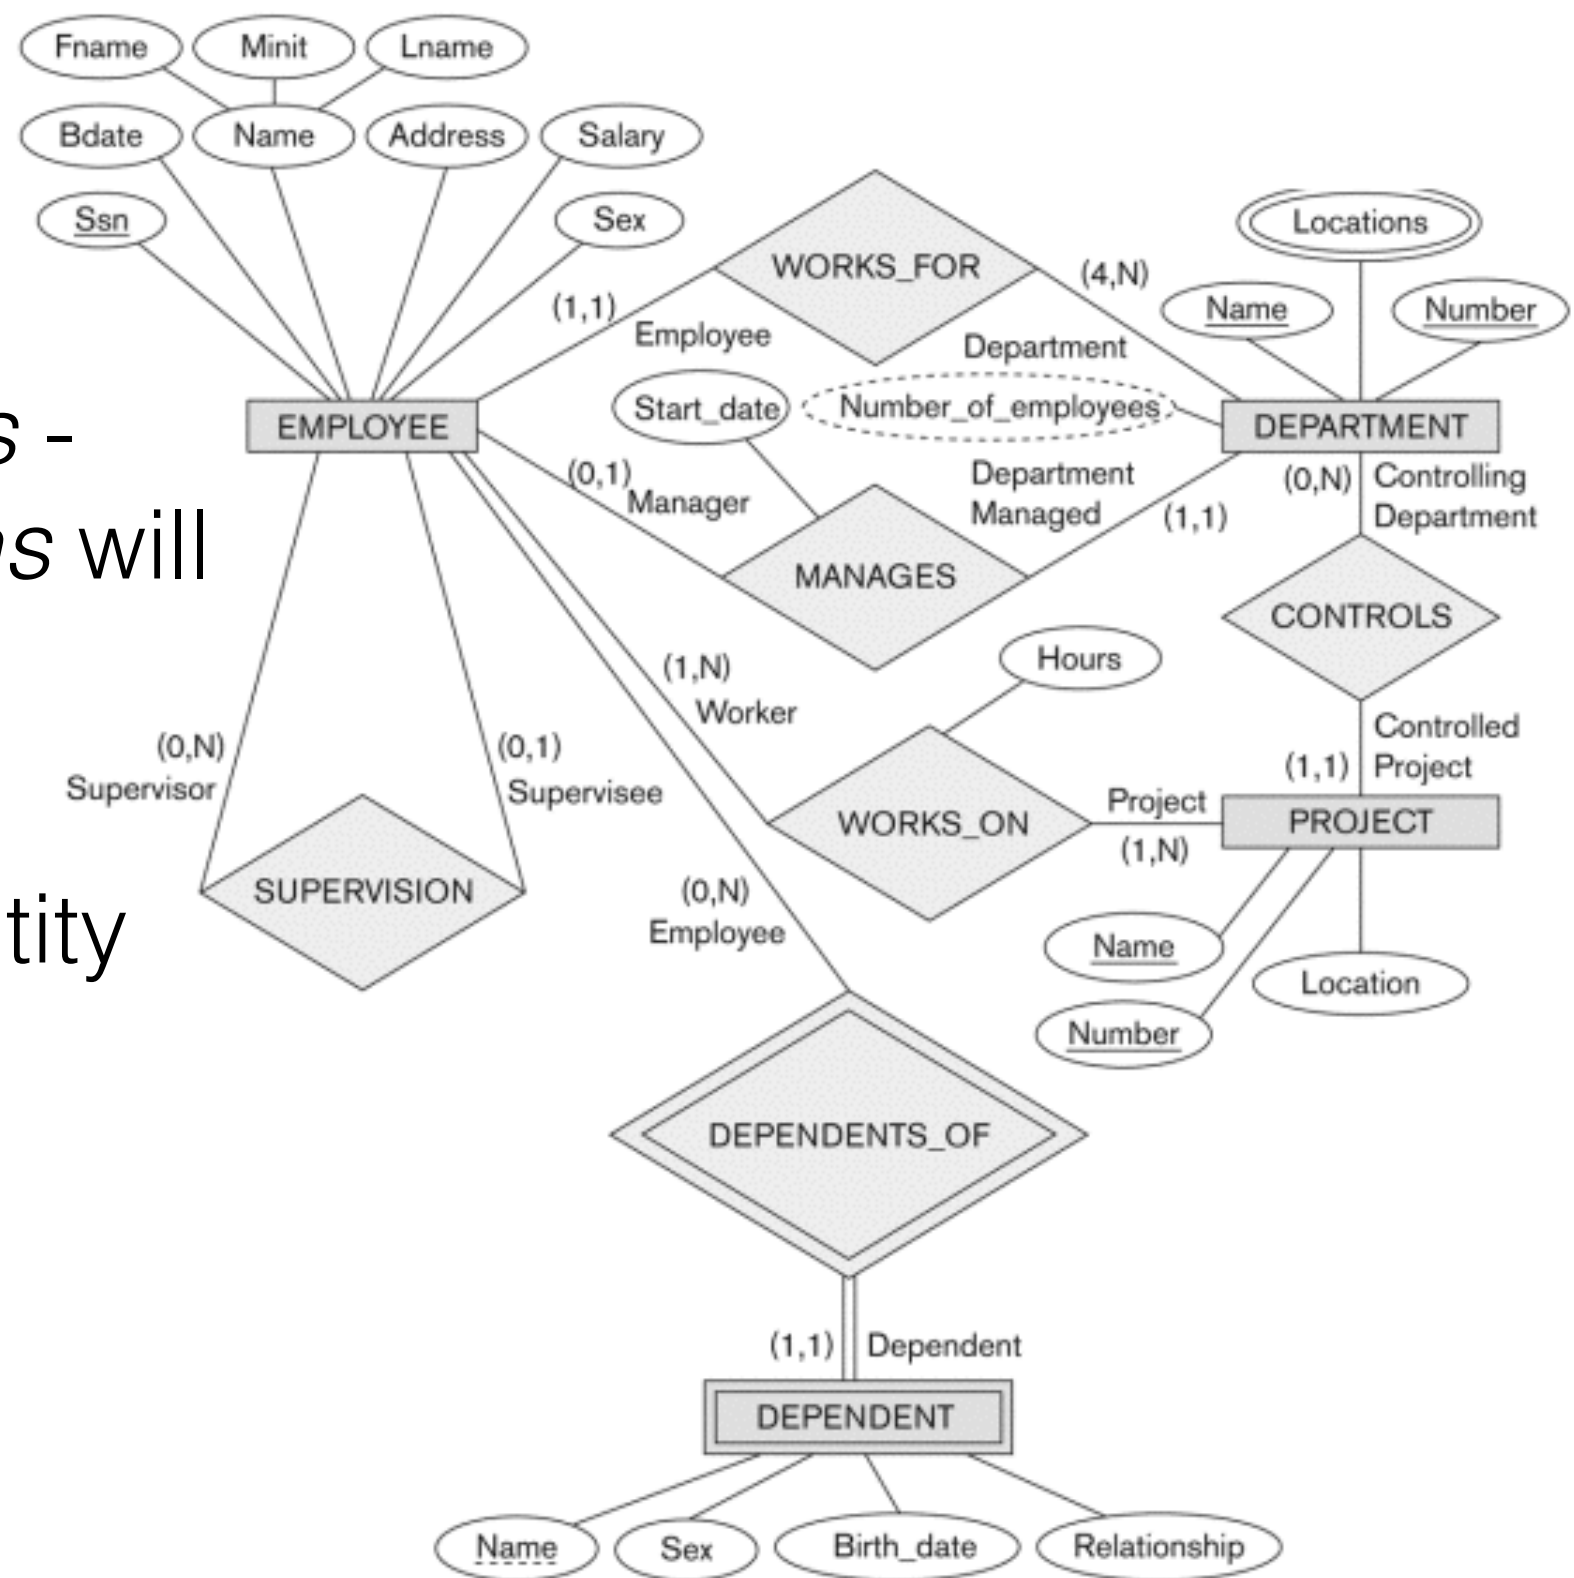
Administrivia

- **Quiz 1** Was Last Week
- Grades will be posted on KEATs within 2-3 weeks!
- Release QUIZ Solutions end of Term (after MCFs)
- **Quiz 2** (SQL) NEXT WEEK DURING LECTURE

Review: ER Diagram Concepts

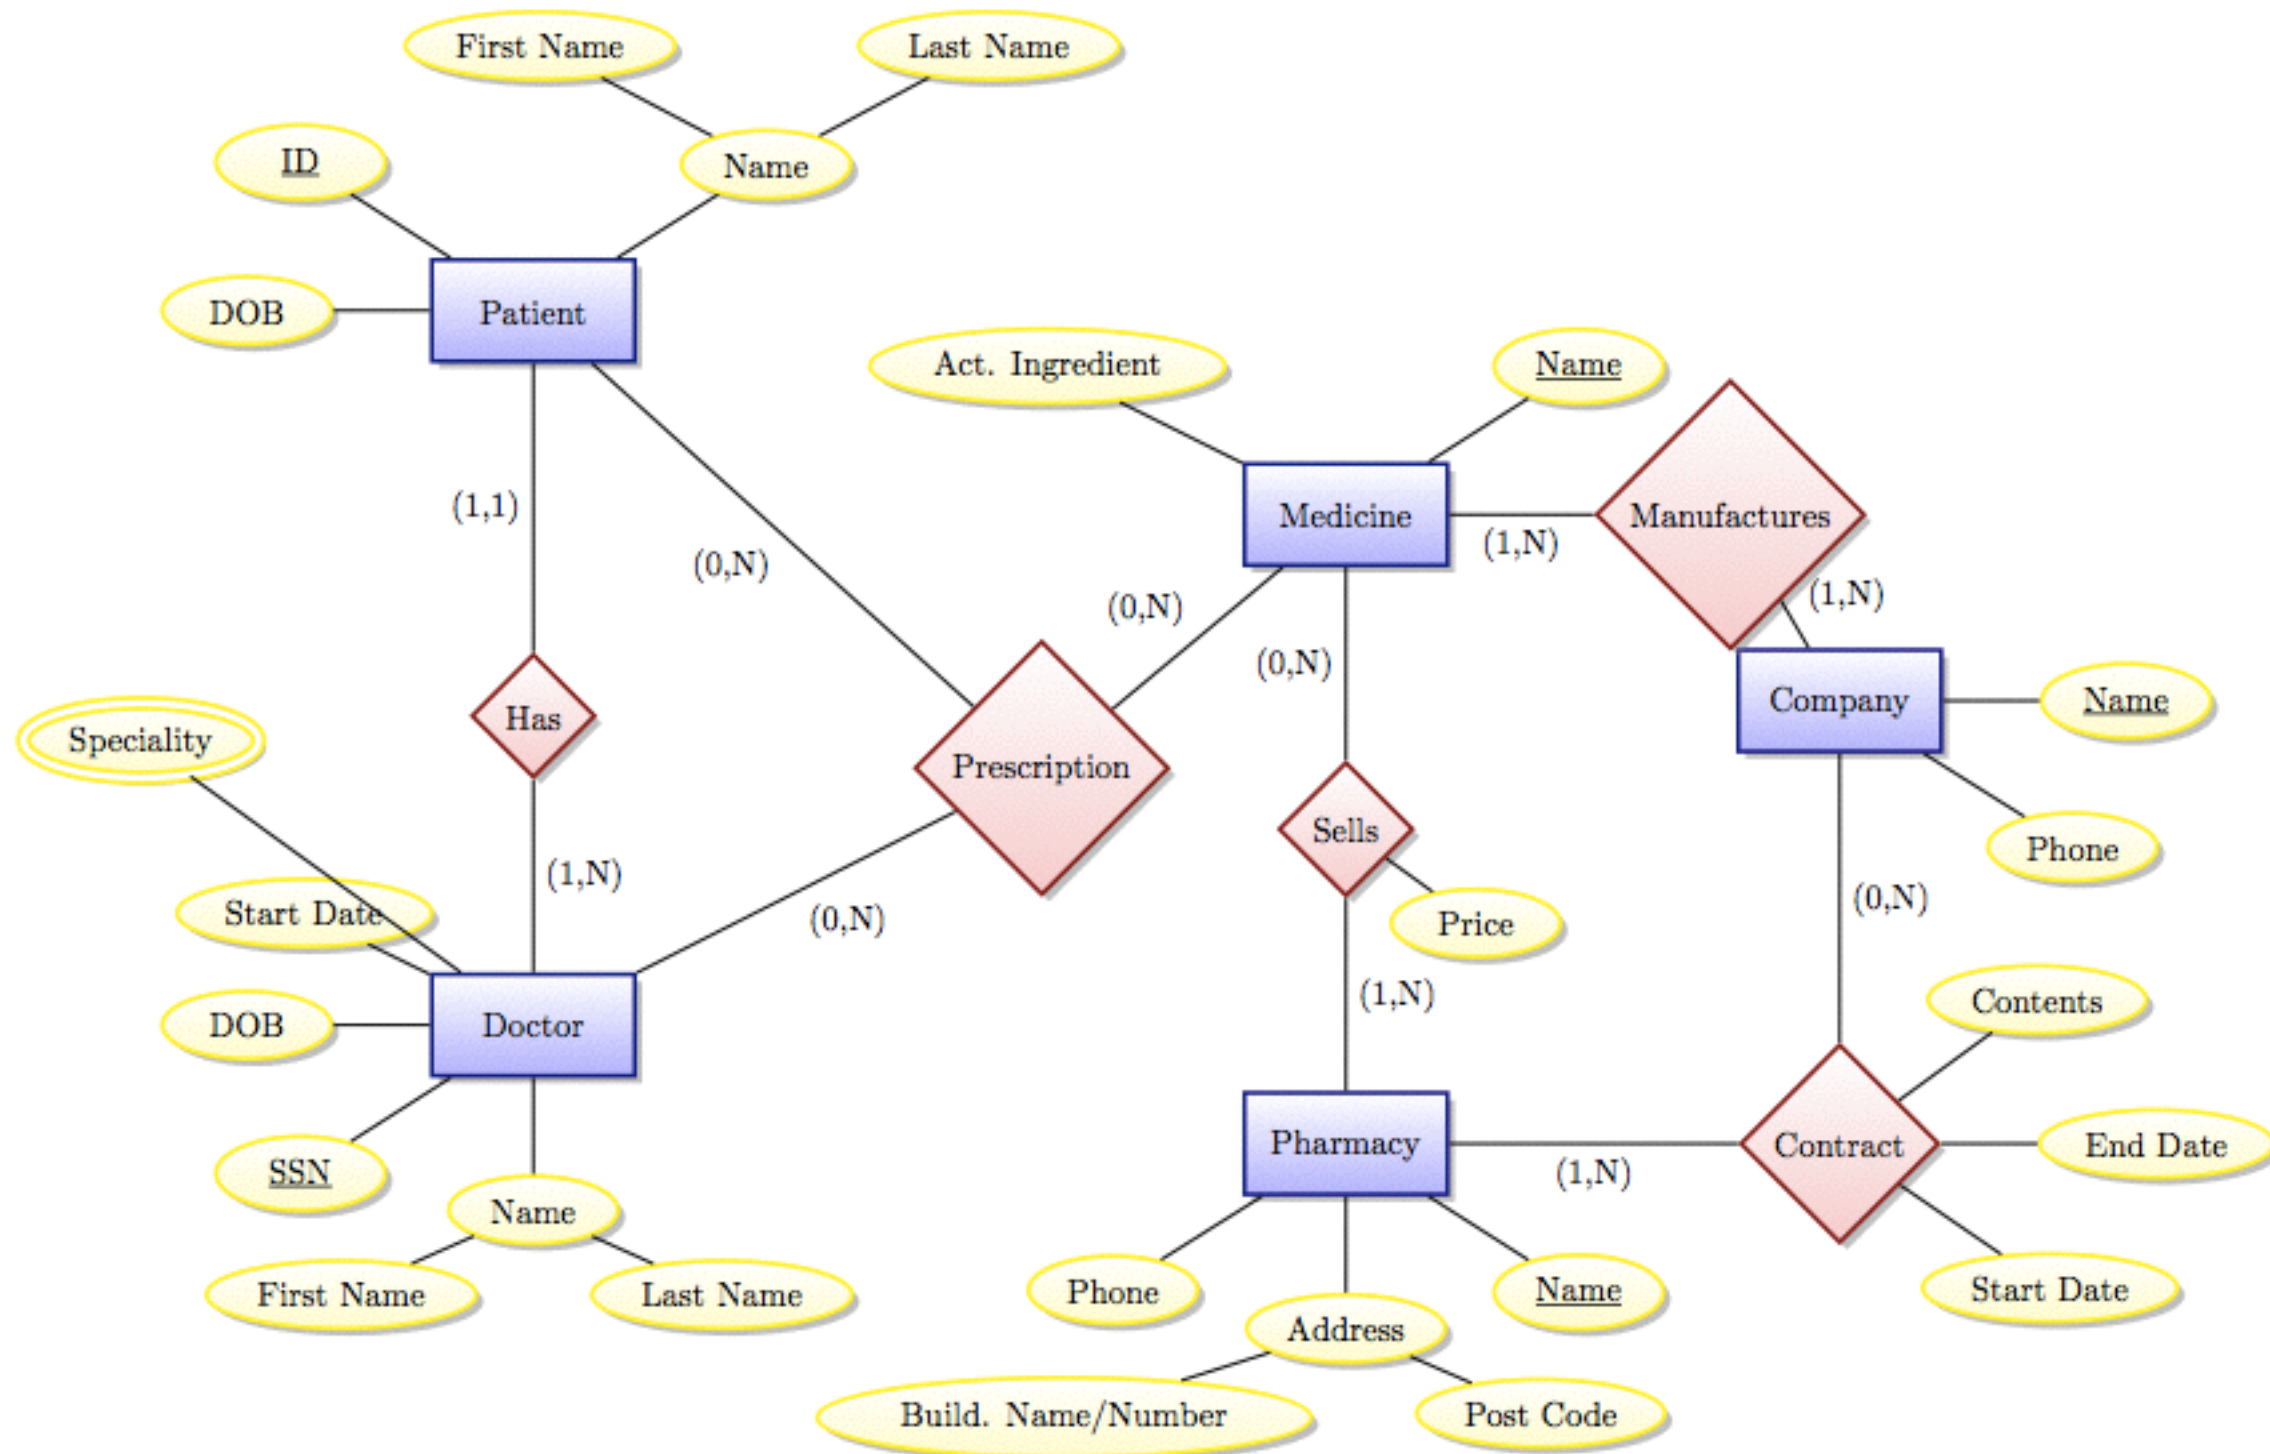
- **Weak Entities**

- in *ER diagrams* - no identifying attributes
- in *Relational Schemas* - Weak Entity's *Relations* will have a PK:
strong entities PK + partial key of weak entity attributes



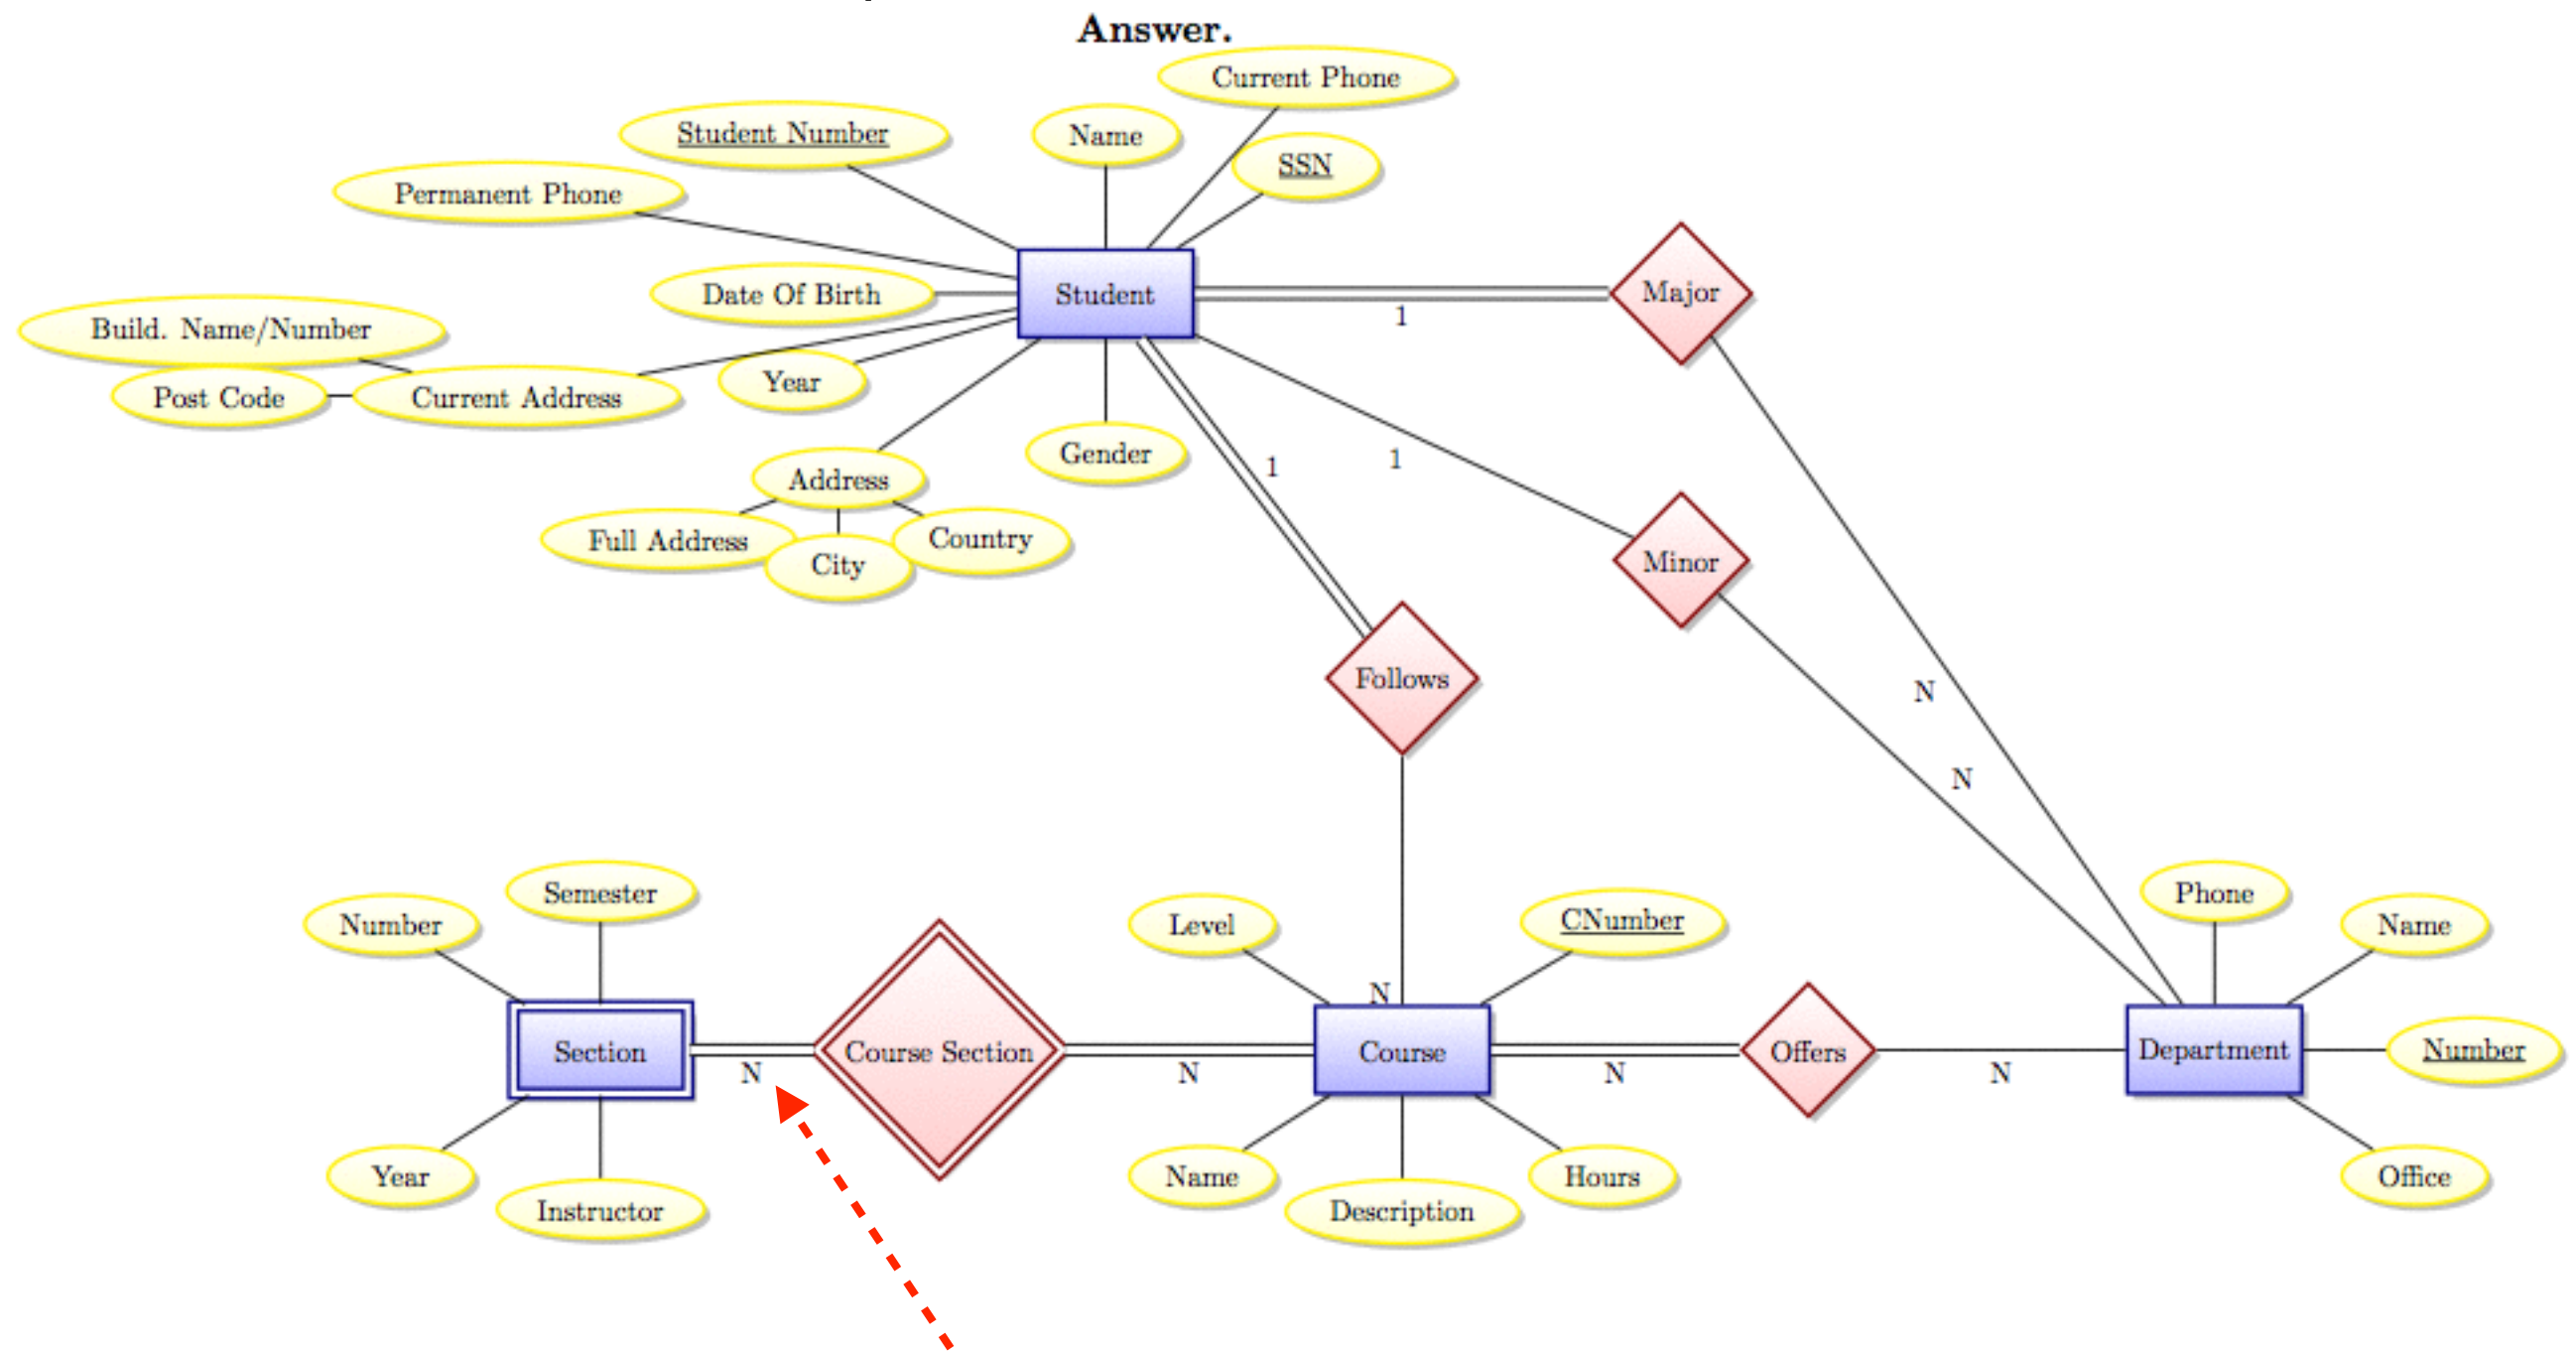
Review: ER Diagram Concepts

- **Multivalued** attribute vs **Composite** attribute



Review: ER Diagram Concepts

- **Total participation** - *entire* set of entities must participate with the relationship



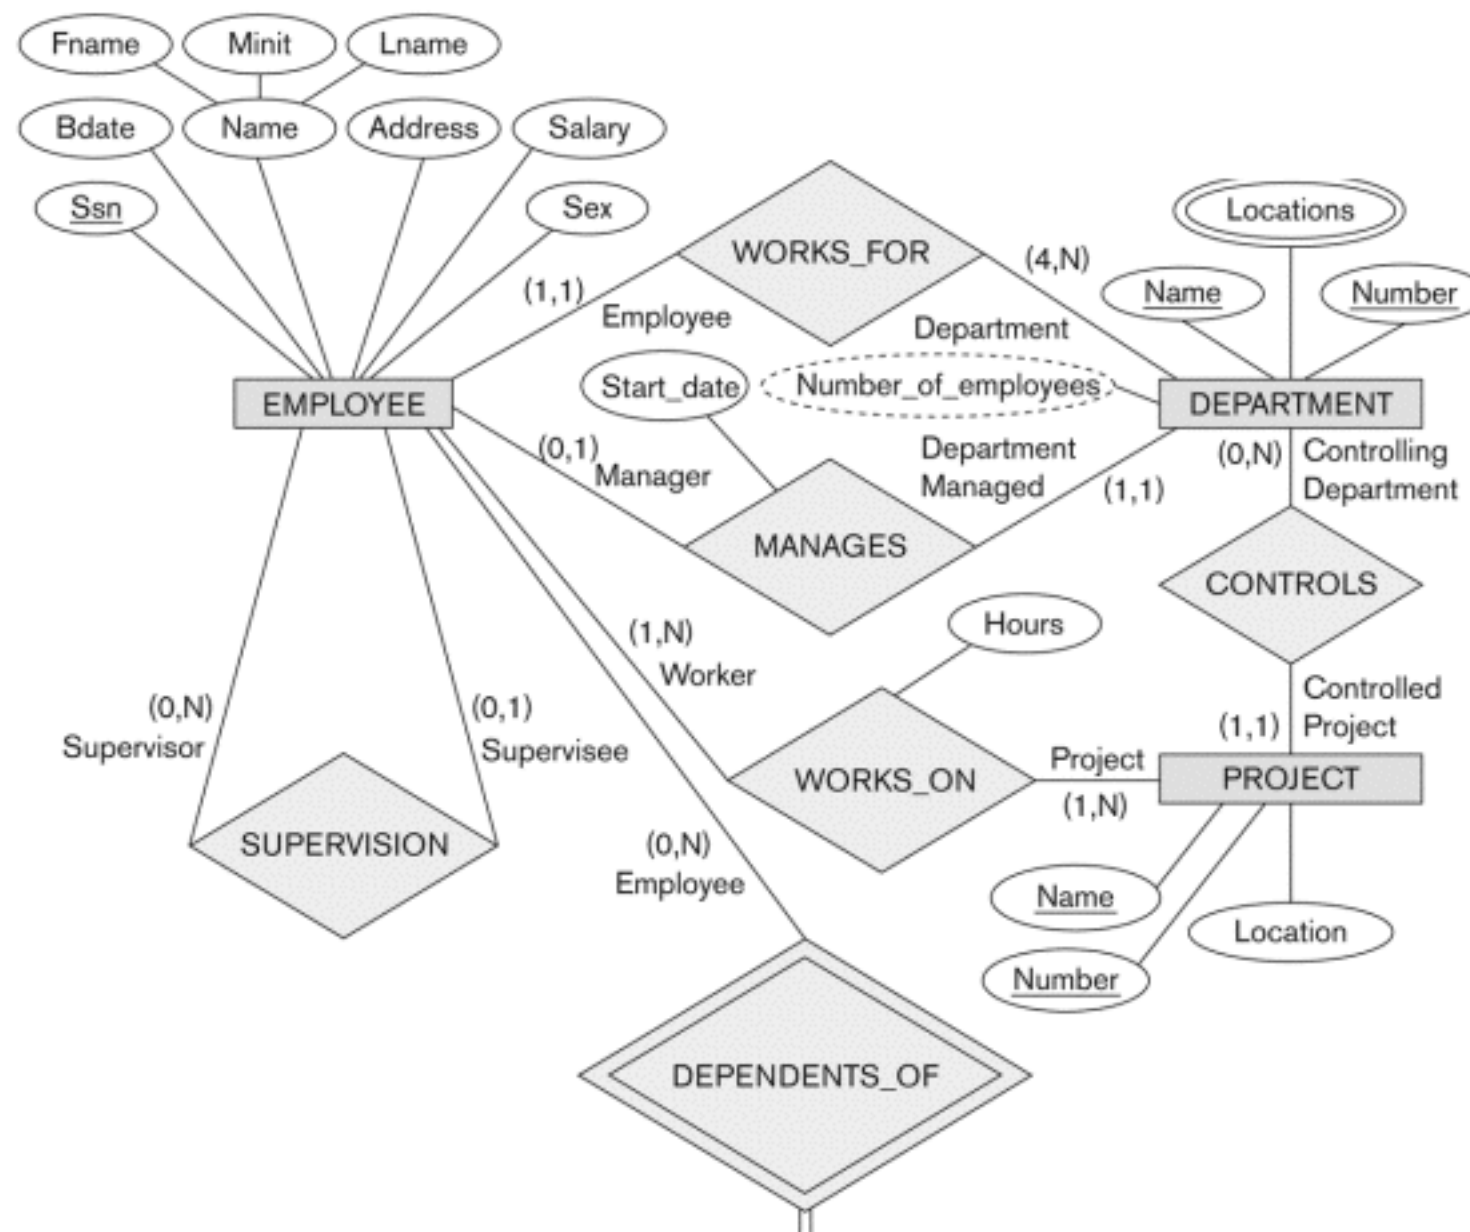
Lab 1 - Part 2 - TYPO!

Review: ER Diagram Concepts

- **Derived attribute -**

not stored in the database, but *derived* via queries / business logic

why? redundancy, and danger of inconstancy



Recap: SQL from Last Week

- Data Definition Language (DDL):

CREATE SCHEMA <schema name>

CREATE TABLE (<attribute definition list>)

ALTER TABLE <table name> **ADD** <attribute definition>

DROP TABLE <table name>

DROP SCHEMA <schema name>

- Data Manipulation Language (DML):

SELECT <attribute list>

FROM <table list>

[WHERE <condition>]

Today

- **DML (Data Manipulation Language) Commands**
 - `SELECT` SQL Queries (continue with `SELECT`)
 - Review Set Operations
 - Math Expressions
 - Casting Types
 - `INSERT` Data
 - `UPDATE` Data
 - `DELETE` Data
- Nested Queries
- Types of Joins
- Grouping and Aggregation Functions
- Assertions and Views

COMPANY Relational Database Schema

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

DEPENDENT

<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
-------------	-----------------------	-----	-------	--------------

COMPANY Populated Database

EMPLOYEE	FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
	Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
	James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1


					DEPT_LOCATIONS	<u>DNUMBER</u>	<u>DLOCATION</u>
						1	Houston
						4	Stafford
						5	Bellaire
						5	Sugarland
						5	Houston
DEPARTMENT	DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE			
	Research	5	333445555	1988-05-22			
	Administration	4	987654321	1995-01-01			
	Headquarters	1	888665555	1981-06-19			

WORKS_ON	<u>ESSN</u>	<u>PNO</u>	HOURS
	123456789	1	32.5
	123456789	2	7.5
	666884444	3	40.0
	453453453	1	20.0
	453453453	2	20.0
	333445555	2	10.0
	333445555	3	10.0
	333445555	10	10.0
	333445555	20	10.0
	999887777	30	30.0
	999887777	10	10.0
	987987987	10	35.0
	987987987	30	5.0
	987654321	30	20.0
	987654321	20	15.0
	888665555	20	null

PROJECT	PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
	ProductX	1	Bellaire	5
	ProductY	2	Sugarland	5
	ProductZ	3	Houston	5
	Computerization	10	Stafford	4
	Reorganization	20	Houston	1
	Newbenefits	30	Stafford	4

DEPENDENT	<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
	333445555	Alice	F	1986-04-05	DAUGHTER
	333445555	Theodore	M	1983-10-25	SON
	333445555	Joy	F	1958-05-03	SPOUSE
	987654321	Abner	M	1942-02-28	SPOUSE
	123456789	Michael	M	1988-01-04	SON
	123456789	Alice	F	1988-12-30	DAUGHTER
	123456789	Elizabeth	F	1967-05-05	SPOUSE

No NEED to RUN DB at HOME ...

SQL Fiddle

MySQL 5.6

View Sample Fiddle

Clear

Text to DDL

User Options

Donate

Flattr

About

1

-- COMPANY Database

2

3

CREATE TABLE EMPLOYEE (FNAME VARCHAR(64), MINIT CHAR, LNAME VARCHAR(64),

4

5

CREATE TABLE DEPARTMENT(DNAME VARCHAR(16) NOT NULL, DNUMBER INTEGER NOT NULL,

6

7

CREATE TABLE DEPT_LOCATIONS(DNUMBER INTEGER NOT NULL, DLOCATION VARCHAR(16) NOT NULL,

8

9

CREATE TABLE PROJECT(PNAME VARCHAR(64) NOT NULL, PNUMBER INTEGER NOT NULL,

10

11

CREATE TABLE WORKS_ON(ESSN CHAR(9), PNO INTEGER NOT NULL, HOURS REAL NOT NULL,

12

13

CREATE TABLE DEPENDENT(ESSN CHAR(9) NOT NULL, DEPENDENT_NAME VARCHAR(64) NOT NULL,

14

15

-- INSERT Some Data

16

-- EMPLOYEE

17

INSERT INTO EMPLOYEE VALUES ("John", "B", "Smith", "123456789", "123456789",

18

INSERT INTO EMPLOYEE VALUES ("Franklin", "T", "Wong", "333445555", "333445555",

Build Schema

Edit Fullscreen

Browser

[:]

1

Run SQL

Edit Fullscreen

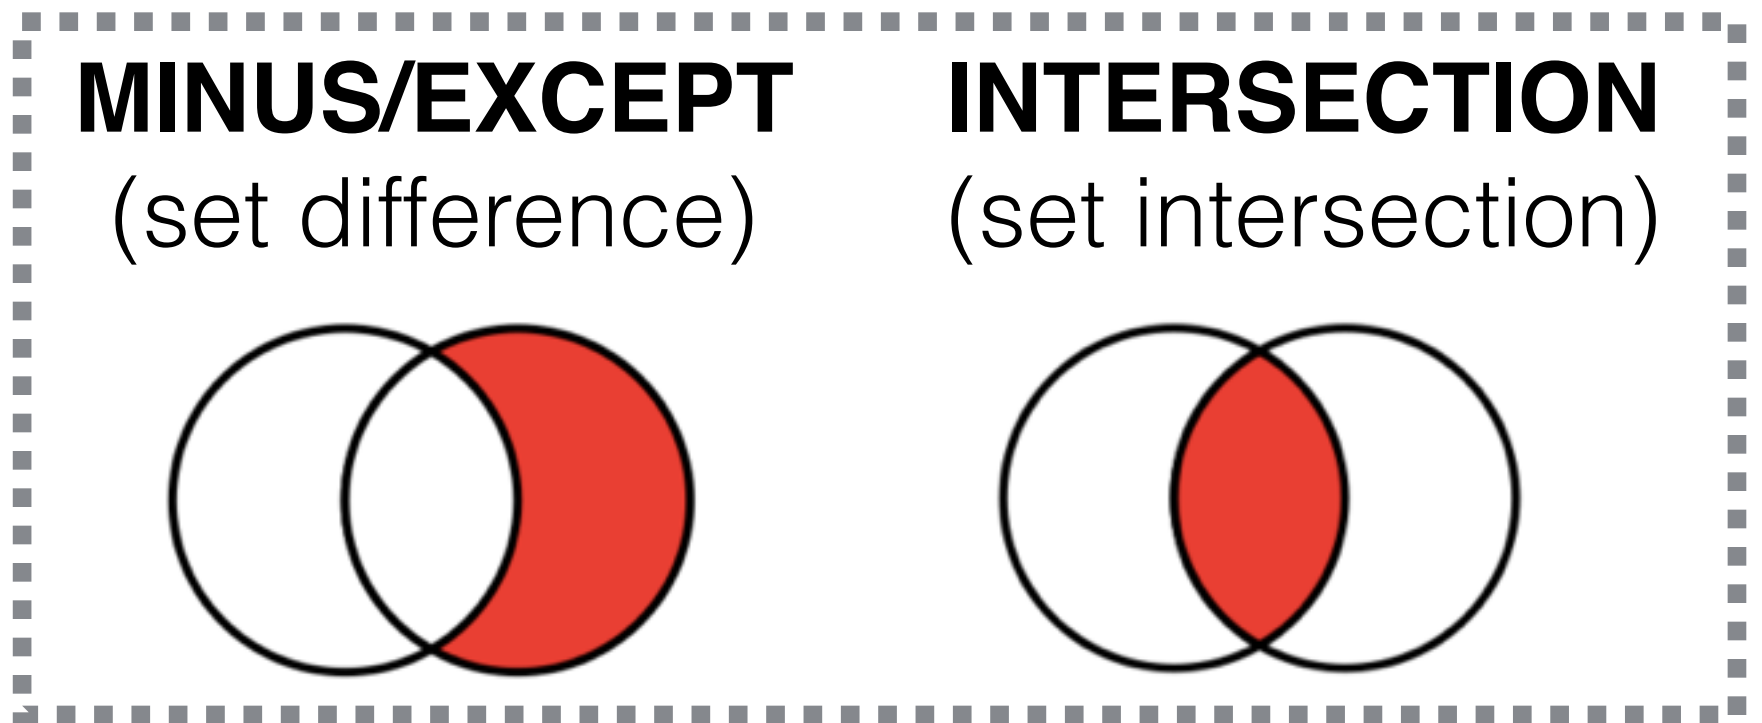

Format Code

[:]

COMPANY DB: <http://sqlfiddle.com/#!9/78de75>

Review: Set Operations

- SQL has directly incorporated some set operations



- Resulting relations of these set operations are sets of tuples — *duplicate tuples are eliminated from the result*
- Set operations apply only to *union compatible relations*:
 - Two relations must have the same attributes (names)
 - Attributes must appear in the same order

Review: Set Operations — UNION

Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
( SELECT      PNUMBER
  FROM        PROJECT, DEPARTMENT, EMPLOYEE
  WHERE       DNUM=DNUMBER AND MGRSSN=SSN
              AND LNAME='Smith' )

UNION

( SELECT      PNUMBER
  FROM        PROJECT, WORKS_ON, EMPLOYEE
  WHERE       PNUMBER=PNO AND ESSN=SSN
              AND LNAME='Smith' )
```

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

WORKS_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

Recap: Set Operations — UNION

Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
( SELECT      PNUMBER
  FROM      PROJECT, DEPARTMENT, EMPLOYEE
  WHERE      DNUM=DNUMBER AND MGRSSN=SSN
            AND LNAME='Smith' )
```

empty

```
UNION
( SELECT      PNUMBER
  FROM      PROJECT, WORKS_ON, EMPLOYEE
  WHERE      PNUMBER=PNO AND ESSN=SSN
            AND LNAME='Smith' )
```

PNUMBER
1
2

EMPLOYEE									
FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO

WORKS_ON		
<u>ESSN</u>	<u>PNO</u>	HOURS

DEPARTMENT			
DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE

PROJECT			
PNAME	<u>PNUMBER</u>	PLOCATION	DNUM

Set Operations — EXCEPT (MINUS)

Example: List SSNs from all employees except those who are working on Project 1.

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

WORKS_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

```
( SELECT      SSN
  FROM      EMPLOYEE )
EXCEPT
( SELECT      ESSN as SSN
  FROM      WORKS_ON
  WHERE      PNO=1 )
```

Set Operations — EXCEPT (MINUS)

Example: List SSNs from all employees except those who are working on Project 1.

(SELECT
FROM

EXCEPT

(SELECT

FROM

WHERE

SSN

EMPLOYEE)

ESSN as SSN

WORKS_ON

PNO=1)

SSN
123456789
333445555
453453453
666884444
888665555
987654321
987987987
999887777

EXCEPT

ESSN
123456789
453453453

SSN
333445555
666884444
888665555
987654321
987987987
999887777

Arithmetic Operations

- The standard arithmetic operators **'+', '-', '*', and '/'** (for addition, subtraction, multiplication, and division, respectively) can be applied to numeric values in an SQL query result

Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
SELECT FNAME, LNAME, 1.1*SALARY  
FROM   EMPLOYEE, WORKS_ON, PROJECT  
WHERE  SSN=ESSN AND PNO=PNUMBER AND  
       PNAME='ProductX'
```


More on Arithmetic Expressions and IF() function

- Constants are allowed
- Note the use of **AS** to alias the results as an attribute
- **IF**(<condition>, <True Value>, <False Value>)
- **IF** with more than 2 values? See **CASE()**, or nested IF statements

```
SELECT  FNAME, LNAME,  
        (SALARY / 1000) AS SALARY_K,  
        1 AS ONE,  
        IF(SALARY > 30000, True, False) AS IS_LOADED,  
        IF(SUPERSSN IS NULL, "Boss", "Worker") AS ETYPE  
FROM    EMPLOYEE
```

FNAME	LNAME	SALARY_K	ONE	IS_LOADED	ETYPE
John	Smith	33.0000	1	1	Worker
Franklin	Wong	40.0000	1	1	Worker
Joyce	English	25.0000	1	0	Worker
Ramesh	Narayan	38.0000	1	1	Worker
James	Borg	55.0000	1	1	Boss
Jennifer	Wallace	43.0000	1	1	Worker
Ahmad	Jabber	25.0000	1	0	Worker
Alicia	Zelaya	25.0000	1	0	Worker

Use of CAST()

- Convert the Data Type of an attribute using **CAST()**:
- **CAST**(<expression> **AS** <type>):

- `BINARY[(N)]`

- `CHAR[(N)]`

- `DATE`

- `DATETIME`

- `DECIMAL[(M[, D])]`

- `SIGNED [INTEGER]`

- `TIME`

- `UNSIGNED [INTEGER]`

```
SELECT FNAME, LNAME,  
       CAST( (SALARY / 1000) AS UNSIGNED)  
       AS SALARY_K  
FROM   EMPLOYEE
```

fname	lname	SALARY_K
John	Smith	30
Franklin	Wong	40
Joyce	English	25
Ramesh	Narayan	38
James	Borg	55
Jennifer	Wallace	43
Ahmad	Jabber	25
Alicia	Zelaya	25

Update: Attribute Data Types - Numerics - Real Numbers

- **Approximate Value** (remember CS1?):
 - FLOAT, REAL, DOUBLE
 - Can specify digit precision (DB does *rounding*)
 - FLOAT(**n**)
 - n - precision, as in ***number of bits used to store the mantissa of the float number in scientific notation***
 - *Example: FLOAT(24) holds a single precision floating point number*
- **Exact Value** (Fixed-Point Type)
 - DECIMAL(i,j) for exact formatted numbers
 - i - precision, ***total number of digits to store number***
 - j - scale, *after decimal*
 - *Example: DECIMAL(7,4) will look like -999.9999 when displayed*
 - More expensive to store, but more precise

Manipulating Data in SQL

- Three SQL commands to modify the STATE of a database (part of the DML for SQL)
 1. **INSERT**
 2. **DELETE**
 3. **UPDATE**
- *They do not modify the SCHEMA of the database*
 - *What commands are used in this case?*
- *Note that **SELECT** is widely considered part of the DML because it clearly is not a DDL command.*

INSERT

- In its simplest form, it is used to add one or more tuples to a relation

```
INSERT INTO  <table name>  
VALUES      <tuple>
```

- Attribute values should be listed in *the same order* as the attributes were specified in the **CREATE TABLE** command

```
INSERT INTO  EMPLOYEE  
VALUES ( 'Richard', 'K', 'Marini', '653298653',  
          '30-DEC-52', '98 Oak Forest, Katy, TX',  
          'M', 37000, '987654321', 4 )
```

INSERT (Specify Values)

- Alternate form of INSERT specifies explicitly the attribute names that correspond to the values in the new tuple
- *Left out attributes will be default value or NULL*

```
INSERT INTO  <table name> (<attribute list>)  
VALUES      <tuple>
```

Example: Insert a tuple for a new EMPLOYEE for whom we only know the FNAME, LNAME, and SSN attributes.

```
INSERT INTO EMPLOYEE (FNAME, LNAME, SSN)  
VALUES ('Richard', 'Marini', '653298653')
```


INSERT Multiple Values from a CREATE

- Another variation of INSERT allows insertion of *multiple tuples* resulting from a query into a relation

Example: We want to create a temporary table that has the employee last name, project name, and hours per week for each employee working on a project.

- *First create a table, WORKS_ON_INFO*

```
CREATE TABLE WORKS_ON_INFO (  
    EMP_NAME          VARCHAR(15) ,  
    PROJ_NAME         VARCHAR(15) ,  
    HOURS_PER_WEEK    DECIMAL(3,1)  
);
```

INSERT Multiple Values from a CREATE

Example: We want to create a temporary table that has the employee last name, project name, and hours per week for each employee working on a project.

- *Then load WORKS_ON_INFO with the results of a joined query:*

```
INSERT INTO WORKS_ON_INFO(EMP_NAME, PROJ_NAME, HOURS_PER_WEEK)
SELECT      E.Lname, P.Pname, W.Hours
FROM        PROJECT P, WORKS_ON W, EMPLOYEE E
WHERE        P.Pnumber=W.Pno AND W.Essn = E.Ssn;
```

INSERT Multiple Values from a CREATE

Example: We want to create a temporary table that has the employee last name, project name, and hours per week for each employee working on a project.

- *Then load WORKS_ON_INFO with the results of a joined query:*

```
INSERT INTO WORKS_ON_INFO(EMP_NAME, PROJ_NAME, HOURS_PER_WEEK)
SELECT      E.Lname, P.Pname, W.Hours
FROM        PROJECT P, WORKS_ON W, EMPLOYEE E
WHERE       P.Pnumber=W.Pno AND W.Essn = E.Ssn;
```

- *Values are mapped to Attributes in the Order they appear*

Using CREATE TABLE ... AS

- It is also possible to do the previous two queries in one CREATE TABLE command (note that **AS** is used in two different ways in the following query):

```
CREATE TABLE WORKS_ON_INFO AS
SELECT      E.Lname AS Emp_Name,
            P.Pname AS Proj_Name,
            W.Hours AS Hours_per_week
FROM        PROJECT P, WORKS_ON W, EMPLOYEE E
WHERE       P.Pnumber=W.Pno AND W.Essn = E.Ssn;
```

- *Note the use of the keyword **AS** to specify table/attribute names*

DELETE

- Removes tuples from a relation

DELETE FROM <table name>
WHERE <condition>

- Includes a WHERE-clause to select the tuples to be deleted
- Referential integrity should be enforced
- Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint)
- A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table
- The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause

DELETE Examples

Which tuples do these queries delete?

- **Example 1** `DELETE FROM EMPLOYEE
WHERE LNAME='Brown';`
- **Example 2** `DELETE FROM EMPLOYEE
WHERE SSN='123456789';`
- **Example 3** `DELETE FROM EMPLOYEE
WHERE DNO = 5;`
- **Example 4** `DELETE FROM EMPLOYEE;`

UPDATE

- Used to modify attribute values of one or more selected tuples

```
UPDATE <table name>  
SET      <attribute>=<value>,...  
WHERE    <condition>
```

- A WHERE-clause selects the tuples to be modified
- An additional SET-clause specifies the attributes to be modified and their new values
- Each command modifies tuples *in the same relation*
- Referential integrity should be enforced

UPDATE Example

Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
UPDATE PROJECT
SET      PLOCATION = 'Bellaire',
        DNUM = 5
WHERE    PNUMBER=10;
```

UPDATE Example

Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
UPDATE PROJECT  
SET      PLOCATION = 'Bellaire',  
         DNUM = 5;  
WHERE PNUMBER=10;
```

UPDATE Example 2

Example: Give all employees in the 'Research' department a 10% raise in salary.

```
UPDATE  EMPLOYEE
SET      SALARY = SALARY*1.1
WHERE    DNO IN (SELECT DNUMBER
                  FROM DEPARTMENT
                  WHERE DNAME='Research' );
```

- Math Expression: the modified SALARY value depends on the original SALARY value in each tuple.
 - The reference to the SALARY attribute on the **right** of = refers to the **old** SALARY value before modification
 - The reference to the SALARY attribute on the **left** of = refers to the **new** SALARY value after modification

Nesting of Queries

- A complete SELECT query, called a *nested query*, can be specified within the WHERE-clause of another query, called the *outer query*
 - Many of the previous queries can be specified in an alternative form using nesting

Example: Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT  FNAME, LNAME, ADDRESS
FROM    EMPLOYEE, DEPARTMENT
WHERE   DNAME= 'Research' AND DNUMBER=DNO
```

Nesting of Queries — Equivalent Query using IN

- Using the comparison operator **IN**:
Compares a value v with a set of values V and returns TRUE if v is one of the elements in V .

Example: Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT  FNAME, LNAME, ADDRESS
FROM    EMPLOYEE
WHERE   DNO IN (SELECT DNUMBER
                FROM DEPARTMENT
                WHERE DNAME='Research' );
```

Nesting of Queries — Explanation of IN

- The nested query selects the number of the 'Research' department
- The outer query select an EMPLOYEE tuple if its DNO value is in the result of either nested query
- **The comparison operator IN compares a value v with a set (or multi-set) of values V, and evaluates to TRUE if v is one of the elements in V**
- In general, we can have several levels of nested queries
- A reference to an *unqualified attribute* refers to the relation declared in the *innermost nested query*
- In previous example, the nested query is *not correlated* with the outer query

Correlated Nested Queries

- If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query*, the two queries are said to be **correlated**
- The result of a correlated nested query is **different** for each tuple (or combination of tuples) of the relation(s) in the outer query.

Example: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT  E.FNAME, E.LNAME
FROM    EMPLOYEE AS E
WHERE   E.SSN IN (SELECT ESSN
                  FROM    DEPENDENT
                  WHERE   ESSN=E.SSN AND
                          E.FNAME=DEPENDENT_NAME )
```


Correlated Nested Queries - Re-written

- A query written with nested SELECT... FROM... WHERE... blocks and using the = or IN comparison operators can ***always*** be expressed as a single block query.

```
SELECT      E.FNAME, E.LNAME
FROM        EMPLOYEE AS E
WHERE       E.SSN IN (SELECT  ESSN
                        FROM    DEPENDENT
                        WHERE    ESSN=E.SSN AND
                                E.FNAME=DEPENDENT_NAME)
```

is re-written as a single block query:

```
SELECT      E.FNAME, E.LNAME
FROM        EMPLOYEE E, DEPENDENT D
WHERE       E.SSN=D.ESSN AND
            E.FNAME=D.DEPENDENT_NAME
```

The EXISTS Function

- EXISTS is used to check whether the result of a correlated nested query is empty (contains no tuples) or not

Example: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
SELECT FNAME, LNAME
FROM   EMPLOYEE E
WHERE  EXISTS (SELECT *
                FROM   DEPENDENT
                WHERE   E.SSN=ESSN AND
                        E.FNAME=DEPENDENT_NAME)
```

Again re-written in a different form...

The EXISTS Function - NOT EXISTS

- NOT EXISTS is TRUE if there are NO tuples as a result of the query.

Example: Retrieve the names of employees who have no dependents.

```
SELECT FNAME, LNAME
FROM   EMPLOYEE E
WHERE  NOT EXISTS (SELECT *
                    FROM   DEPENDENT
                    WHERE  E.SSN=ESSN)
```

ALL comparison operator

- Comparison operators to compare a *single* value (as an attribute) to a *set* or *multiset* (a nested query)

Example: Retrieve the names of employees whose salary is greater than the salary of all employees in department 5.

```
SELECT      LNAME, FNAME
FROM        EMPLOYEE
WHERE       SALARY > ALL( SELECT  SALARY
                           FROM    EMPLOYEE
                           WHERE   DNO=5 )
```

NULLs in SQL Queries

- SQL allows queries that check if a value is **NULL** (missing or undefined or not applicable)
- SQL uses **IS** or **IS NOT** to compare NULLs because it considers each NULL value distinct from other NULL values, so *equality comparison is not appropriate*
- In join conditions, tuples with NULL values in these attributes are not included in result (i.e. DNUMBER = DNO, and both are NULL)

Example: Retrieve the names of all employees who do not have supervisors

```
SELECT    FNAME, LNAME  
FROM      EMPLOYEE  
WHERE     SUPERSSN IS NULL
```

BREAKTIME

Joined Relations — Using JOIN

- Using the JOIN keyword can specify “joined relations”
- Two joined relations look like any other relation
- Many types:
 - **JOIN** (regular “theta” join as you will see in Rel. Alg.)
 - **NATURAL JOIN**
 - **LEFT OUTER JOIN, LEFT JOIN**
 - **RIGHT OUTER JOIN, RIGHT JOIN**
 - **FULL OUTER JOIN, OUTER JOIN**
 - **INNER JOIN**
 - **CROSS JOIN**

Joined Relations — JOIN ... ON

- SELECT ... with a JOIN Condition in the WHERE clause:

```
SELECT FNAME, LNAME, ADDRESS
FROM    EMPLOYEE, DEPARTMENT
WHERE   DNAME='Research' AND DNUMBER=DNO
```

- Using JOIN ... ON as an “*equi-join*”

```
SELECT FNAME, LNAME, ADDRESS
FROM    (EMPLOYEE JOIN DEPARTMENT
          ON DNUMBER=DNO)
WHERE   DNAME='Research'
```

- NATURAL JOIN (Note the use of Aliasing with AS)

```
SELECT      FNAME, LNAME, ADDRESS
FROM        (EMPLOYEE NATURAL JOIN DEPARTMENT
              AS DEPT(DNAME, DNO, MSSN, MSDATE))
WHERE       DNAME='Research'
```


Distinguishing between JOIN Functions

- **NATURAL JOIN** (same as JOIN):
 - *No join condition may be specified, implicit condition to join on attributes with the same name*
- **INNER JOIN:**
 - *Tuple is included in the result only if a matching tuple exists in the other relation (default type of JOIN)*
- **OUTER JOIN:**
 - *all matching tuples are returned (depending on type of OUTER JOIN):*
 - LEFT OUTER JOIN
 - RIGHT OUTER JOIN
 - FULL OUTER JOIN

OUTER JOIN Examples

- Previous query:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   EMPLOYEE E S  
WHERE  E.SUPERSSN=S.SSN
```

- How does changing it to a **LEFT OUTER JOIN** modify the results?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   (EMPLOYEE E LEFT OUTER JOIN  
        EMPLOYEE AS S  
        ON E.SUPERSSN=S.SSN)
```

OUTER JOIN Examples

- Previous query:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   EMPLOYEE E S  
WHERE  E.SUPERSSN=S.SSN
```

- How does changing it to a **LEFT OUTER JOIN** modify the results?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   (EMPLOYEE E LEFT OUTER JOIN  
        EMPLOYEE AS S  
        ON E.SUPERSSN=S.SSN)
```

Same as above but include this row...

James	Borg	NULL	NULL
-------	------	------	------

OUTER JOIN Examples

- Previous query:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   EMPLOYEE E S  
WHERE  E.SUPERSSN=S.SSN
```

- How does changing it to a **RIGHT OUTER JOIN** modify the results?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   (EMPLOYEE E RIGHT OUTER JOIN  
        EMPLOYEE AS S  
        ON E.SUPERSSN=S.SSN)
```

OUTER JOIN Examples

- Previous query:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM   EMPLOYEE E S
WHERE  E.SUPERSSN=S.SSN
```

- How does changing it to a **RIGHT OUTER JOIN** modify the results?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM   (EMPLOYEE E RIGHT OUTER JOIN
        EMPLOYEE AS S
        ON E.SUPERSSN=S.SSN)
```

Same as above but include non-matching tuples from RIGHT table:

NULL	NULL	John	Smith
...			

OUTER JOIN Examples

- Previous query:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM   EMPLOYEE E S
WHERE  E.SUPERSSN=S.SSN
```

- How does changing it to a **FULL OUTER JOIN** modify the results?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM   (EMPLOYEE E FULL OUTER JOIN
        EMPLOYEE AS S
        ON E.SUPERSSN=S.SSN)
```

OUTER JOIN Examples

- Previous query:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM   EMPLOYEE E S
WHERE  E.SUPERSSN=S.SSN
```

- How does changing it to a **FULL OUTER JOIN** modify the results?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM   (EMPLOYEE E FULL OUTER JOIN
        EMPLOYEE AS S
        ON E.SUPERSSN=S.SSN)
```

Include both LEFT and RIGHT table's non-matching tuples:

James	Bord	NULL	NULL
NULL	NULL	Joyce	English
...			

CROSS JOIN — (Cartesian Product)

- What does this query return?

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM    (EMPLOYEE E CROSS JOIN EMPLOYEE S)
```

- Equivalent to:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM    EMPLOYEE E S
```

- As well as:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM    EMPLOYEE E JOIN EMPLOYEE S
```


Multiway JOINS

Example: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
SELECT    PNUMBER, DNUM, LNAME, ADDRESS, BDATE
FROM      PROJECT, DEPARTMENT, EMPLOYEE
WHERE     DNUM=DNUMBER AND MGRSSN=SSN
          AND PLOCATION='Stafford'
```


- Is equivalent to specifying a *multiway* join:

```
SELECT    PNUMBER, DNUM, LNAME, ADDRESS, BDATE
FROM      ((PROJECT JOIN DEPARTMENT ON DNUM=DNUMBER)
           JOIN EMPLOYEE ON MGR_SSN=SSN)
WHERE     PLOCATION='Stafford'
```


Viewed as Venn Diagrams

SQL JOINS


Source: jpromero.com




```
SELECT *  
FROM TableA a  
LEFT JOIN TableB b  
ON a.Key = b.Key
```




```
SELECT *  
FROM TableA a  
LEFT JOIN TableB b  
ON a.Key = b.Key  
WHERE b.Key IS NULL
```




```
SELECT *  
FROM TableA a  
FULL OUTER JOIN TableB b  
ON a.Key = b.Key
```




```
SELECT *  
FROM TableA a  
RIGHT JOIN TableB b  
ON a.Key = b.Key
```



```
SELECT *  
FROM TableA a  
FULL OUTER JOIN TableB b  
ON a.Key = b.Key  
WHERE a.Key IS NULL  
OR b.Key IS NULL
```



```
SELECT *  
FROM TableA a  
RIGHT JOIN TableB b  
ON a.Key = b.Key  
WHERE a.Key IS NULL
```



```
SELECT *  
FROM TableA a  
INNER JOIN TableB b  
ON a.Key = b.Key
```

Aggregate Functions

Example: Find the maximum salary, the minimum salary, and the average salary among all employees.

```
SELECT    MAX ( SALARY ) , MIN ( SALARY ) , AVG ( SALARY )  
FROM      EMPLOYEE
```

max(salary)	min(salary)	avg(salary)
55000	25000	35125.0000

- Functions include **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- Some SQL implementations may not allow more than one function in the SELECT-clause

Example: Aggregate Functions - postgresql

Example of Statistics Specific - useful in Online Analytics Processing (OLAP)

Table 9-44. Aggregate Functions for Statistics

Function	Argument Type	Return Type	Description
<code>corr(Y, X)</code>	double precision	double precision	correlation coefficient
<code>covar_pop(Y, X)</code>	double precision	double precision	population covariance
<code>covar_samp(Y, X)</code>	double precision	double precision	sample covariance
<code>regr_avgx(Y, X)</code>	double precision	double precision	average of the independent variable ($\text{sum}(X)/N$)
<code>regr_avgy(Y, X)</code>	double precision	double precision	average of the dependent variable ($\text{sum}(Y)/N$)
<code>regr_count(Y, X)</code>	double precision	bigint	number of input rows in which both expressions are nonnull
<code>regr_intercept(Y, X)</code>	double precision	double precision	y-intercept of the least-squares-fit linear equation determined by the (X, Y) pairs
<code>regr_r2(Y, X)</code>	double precision	double precision	square of the correlation coefficient
<code>regr_slope(Y, X)</code>	double precision	double precision	slope of the least-squares-fit linear equation determined by the (X, Y) pairs
<code>regr_sxx(Y, X)</code>	double precision	double precision	$\text{sum}(X^2) - \text{sum}(X)^2/N$ ("sum of squares" of the independent variable)
<code>regr_sxy(Y, X)</code>	double precision	double precision	$\text{sum}(X*Y) - \text{sum}(X) * \text{sum}(Y)/N$ ("sum of products" of independent times dependent variable)
<code>regr_syy(Y, X)</code>	double precision	double precision	$\text{sum}(Y^2) - \text{sum}(Y)^2/N$ ("sum of squares" of the dependent variable)
<code>stddev(expression)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	historical alias for <code>stddev_samp</code>
<code>stddev_pop(expression)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	population standard deviation of the input values
<code>stddev_samp(expression)</code>	smallint, int, bigint, real, double precision, or numeric	double precision for floating-point arguments, otherwise numeric	sample standard deviation of the input values

Aggregate Functions

Example: Find the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department.

```
SELECT    MAX ( SALARY ) , MIN ( SALARY ) , AVG ( SALARY )  
FROM      EMPLOYEE , DEPARTMENT  
WHERE     DNO=DNUMBER AND DNAME='Research'
```


Aggregate Functions - COUNT

Example: Retrieve the total number of employees in the company (Q1), and the number of employees in the 'Research' department (Q2).

Q1 : SELECT COUNT (*)
 FROM EMPLOYEE ;

Q2 : SELECT COUNT (*)
 FROM EMPLOYEE, DEPARTMENT
 WHERE DNO=DNUMBER AND DNAME= 'Research' ;

Aggregate Functions - COUNT

Example: Select the names of the all employees who have two or more dependents.

```
SELECT LNAME, FNAME
FROM    EMPLOYEE
WHERE   (SELECT COUNT(*)
        FROM DEPENDENT
        WHERE SSN=ESS) >= 2;
```

*Note that when the result is one attribute and one tuple – that becomes a **SCALAR**.*

Aggregate Functions - COUNT

Example: Count the number of distinct salary values in the Employees table;

```
SELECT    COUNT (DISTINCT SALARY)  
FROM      EMPLOYEE
```

Note that NULL values are not counted as part of the aggregate

Grouping

- In many cases, we want to apply the aggregate functions to *subgroups of tuples* in a relation
- Each subgroup of tuples consists of the set of tuples that have the *same value* for the *grouping attribute(s)*
- The function is applied to each subgroup independently
- SQL has a **GROUP BY**-clause for specifying the grouping attributes, which *must also appear in the SELECT-clause*

SELECT <attribute list, include **grouping attributes**>
FROM <table list>
[WHERE <condition>
GROUP BY <**grouping attributes**>

Grouping with Aggregate Functions

Example: For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
SELECT      DNO, COUNT ( * ), AVG ( SALARY )  
FROM        EMPLOYEE  
GROUP BY    DNO
```

- EMPLOYEE tuples are divided into groups - Each group having the same value for the grouping attribute **DNO**
- The COUNT and AVG functions are applied to each such group of tuples separately
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples
- A join condition can be used in conjunction with grouping

Grouping with Aggregate Functions

Example: For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
SELECT      DNO, COUNT (*), AVG (SALARY)
FROM        EMPLOYEE
GROUP BY    DNO
```

Fname	Minit	Lname	<u>Ssn</u>	...	Salary	Super_ssn	Dno		Dno	Count (*)	Avg (Salary)
John	B	Smith	123456789		30000	333445555	5		5	4	33250
Franklin	T	Wong	333445555		40000	888665555	5		4	3	31000
Ramesh	K	Narayan	666884444		38000	333445555	5		1	1	55000
Joyce	A	English	453453453	...	25000	333445555	5				
Alicia	J	Zelaya	999887777		25000	987654321	4				
Jennifer	S	Wallace	987654321		43000	888665555	4				
Ahmad	V	Jabbar	987987987		25000	987654321	4				
James	E	Bong	888665555		55000	NULL	1				

Result of Q24

Grouping with Aggregate Functions

Example: For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT      PNAME, PNUMBER, COUNT ( * )  
FROM        PROJECT, WORKS_ON  
WHERE       PNUMBER=PNO  
GROUP BY    PNUMBER, PNAME
```

- The grouping and functions are applied ***after*** joining the two relations

Grouping with Aggregate Functions

```
SELECT      PNAME, PNUMBER, COUNT(*)
FROM        PROJECT, WORKS_ON
WHERE       PNUMBER=PNO
GROUP BY    PNUMBER, PNAME
```

Pname	<u>Pnumber</u>	...	<u>Essn</u>	<u>Pno</u>	Hours
ProductX	1	...	123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666884444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10		333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

PNAME	PNUMBER	COUNT(*)
ProductX	1	2
ProductY	2	3
ProductZ	3	3
Computerization	10	3
Reorganization	20	2
Newbenefits	30	3

The HAVING-Clause

- Sometimes we want to retrieve the values of these functions for only those *groups that satisfy certain conditions*
- The **HAVING**-clause is used for specifying a selection condition on groups (rather than on individual tuples)

SELECT	<attribute list, include grouping attributes >
FROM	<table list>
[WHERE	<condition>]
GROUP BY	< grouping attributes >
HAVING	<condition>

The HAVING-Clause

Example: For each project *on which more than two employees work*, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT      PNUMBER, PNAME, COUNT ( * )
FROM        PROJECT, WORKS_ON
WHERE       PNUMBER=PNO
GROUP BY    PNUMBER, PNAME
HAVING      COUNT ( * ) > 2
```

Pname	<u>Pnumber</u>	...	<u>Essn</u>	<u>Pno</u>	Hours
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
Computerization	10		333445555	10	10.0
Computerization	10	...	999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

Pname	Count (*)
ProductY	3
Computerization	3
Reorganization	3
Newbenefits	3

Result of Q26
(Pnumber not shown)

The HAVING-Clause

```
SELECT      PNUMBER, PNAME, COUNT ( * )  
FROM        PROJECT, WORKS_ON  
WHERE       PNUMBER=PNO  
GROUP BY    PNUMBER, PNAME  
HAVING      COUNT ( * ) > 2
```

Pname	<u>Pnumber</u>	...	<u>Essn</u>	<u>Pno</u>	Hours
ProductX	1		123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666884444	3	40.0
ProductZ	3	...	333445555	3	10.0
Computerization	10		333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by
the HAVING condition `COUNT (*) > 2`

The HAVING-Clause (contd.)

Example: Count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for the departments where more than five employees work.

```
SELECT      DNAME, COUNT ( * )  
FROM        DEPARTMENT, EMPLOYEE  
WHERE       DNUMBER=DNO AND SALARY > 40000  
GROUP BY   DNAME  
HAVING      COUNT ( * ) > 5
```

Is this the correct query?

The HAVING-Clause (contd.)

Example: Count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for the departments where more than five employees work.

```
SELECT      DNAME, COUNT ( * )  
FROM        DEPARTMENT, EMPLOYEE  
WHERE       DNUMBER=DNO AND SALARY > 40000  
GROUP BY    DNAME  
HAVING      COUNT ( * ) > 5
```

Incorrect query! It selects only the departments that have more than five employees who each earn more than \$40,000. The query is too limited.

The HAVING-Clause (contd.)

Example: Count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for the departments where more than five employees work.

```
SELECT  DNAME, COUNT ( * )
FROM    DEPARTMENT, EMPLOYEE
WHERE   DNUMBER=DNO AND SALARY > 40000
        AND DNO IN (SELECT  DNO
                     FROM    EMPLOYEE
                     GROUP BY DNO
                     HAVING  COUNT ( * ) > 5 )
```

Use a Nested Correlated Query – this nested query selects the DEPARTMENTS who's number of employees > 5.

Summary of SQL Queries

- A query in SQL can consist of up to six clauses, but only the first two, ***SELECT and FROM, are mandatory***. The clauses are specified in the following order:

SELECT	<attribute list, include any grouping attributes>
FROM	<table list>
[WHERE	<condition>]
[GROUP BY	<grouping attribute(s)>]
[HAVING	<grouping condition>]
[ORDER BY	<attribute list>]

Summary of SQL Queries

- The **SELECT**-clause lists the attributes or functions to be retrieved
- The **FROM**-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The **WHERE**-clause specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- **GROUP BY** specifies grouping attributes
- **HAVING** specifies a condition for selection of groups
- **ORDER BY** specifies an order for displaying the result of a query
- A query is evaluated by:
 1. applying the WHERE-clause
 2. then GROUP BY and HAVING
 3. and finally the SELECT-clause
 4. and then ORDER BY the resulting tuples

Constraints as Assertions

- General constraints: constraints that do not fit in the basic SQL categories
- Useful for ***Schema Assertions*** - Outside the scope of the built-in relational model constraints (primary and unique keys, entity integrity, referential integrity).
- Defines whether the State of the Database is VALID at any given point of time.
- **CREATE ASSERTION**, Components include:
 - A constraint name
 - Followed by a CHECK keyword
 - Followed by a condition clause
- Enforcing the Assertion is up to the Database Implementation – i.e. Rejecting a Query that will violate the CHECK ASSERTION.

CREATE ASSERTION Example

Example: The salary of an employee must not be greater than the salary of the manager of the department that the employee works for

assertion name

```
CREAT ASSERTION SALARY_CONSTRAINT
CHECK (NOT EXISTS (
    SELECT *
    FROM    EMPLOYEE E, EMPLOYEE M, DEPARTMENT D
    WHERE   E.SALARY > M.SALARY AND
           E.DNO=D.NUMBER AND D.MGRSSN=M.SSN) )
```

assertion condition

Views in SQL

- A view is a “*virtual*” table that is derived from other tables
- Two ways they are implemented in implementation:
 - Query modification - *copy and paste queries*
 - View materialization - *short-term physical implementation*
- Limited for UPDATE operations
 - Table not physical stored
- Allows full query operations
- A convenience for expressing certain operations
- Useful for security and authorization
- Prevents redundant storage of data

SQL View Example

Example: A “friendlier” view of WORKS_ON

- SQL command: CREATE VIEW

view name

CREATE VIEW WORKS_ON1 AS

```
SELECT FNAME, LNAME, PNAME, HOURS
FROM   EMPLOYEE, PROJECT, WORKS_ON
WHERE  SSN=ESSN AND PNO=PNUMBER
```

query to specify the contents of view

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------

view name option: specify attribute names

WORKS_ON1 (FIRST_NAME, LAST_NAME, PROJECT, HOURS)

Using a Virtual Table (a View)

Example: A “friendlier” view of WORKS_ON

- We can specify SQL queries on a newly create table (view):

```
SELECT FNAME, LNAME  
FROM WORKS_ON1  
WHERE PNAME= 'ProductX' ;
```

- When no longer needed, a view can be dropped:

```
DROP WORKS_ON1 ;
```

- Dropping a View does NOT modify the data!

Important Take-Aways

- 6 Parts of SELECT Query and Execution Order
 - SELECT... FROM ... WHERE... GROUP BY ... HAVING... ORDER BY
- Modifying the STATE of a Data Base
 - INSERT, UPDATE, DELETE
- Nested Queries in WHERE clause:
 - IN, EXISTS, ALL in conditions, correlated queries
- Types of Joins: INNER, OUTER, NATURAL JOINS
- GROUP BY... Aggregation Functions
- Why and How to use: CREATE ASSERTIONS
- Why and How to use: CREATE VIEW