

Lecture 10: Sorting Algorithms

(Chapter 11, Sections 11.1, 11.2,
11.3 from the book)

Agenda

◆ Sorting


- Sorting Problem
- Merge Sort
- Quick Sort
- Bucket Sort

◆ Sorting Lower Bound

Sorting Problem

- ◆ Sort a given sequence of data items according to their keys.
 - Examples
 - 1) Input: (15, ...), (26, ...), (11, ...), (23, ...), (7, ...), (31, ...), (30, ...)
Output: (7, ...), (11, ...), (15, ...), (23, ...), (26, ...), (30, ...), (31, ...)
 - 2) Input: ("go", ...), ("did", ...), ("me", ...), ("bet", ...), ("kit", ...)
Output: ("bet", ...), ("did", ...), ("go", ...), ("kit", ...), ("me", ...)
- ◆ Sorting is a fundamental application for computers.
- ◆ Sorting is perhaps the most intensively studied and important operation in computer science.
- ◆ An initial sort of the data can significantly enhance the performance of an algorithm.

Merge Sort



Divide-and-Conquer

◆ **Divide-and conquer** is a general algorithm design paradigm:

- **Divide**: divide the input data S in two disjoint subsets S_1 and S_2
- **Recur**: solve the subproblems associated with S_1 and S_2
- **Conquer**: combine the solutions for S_1 and S_2 into a solution for S

Merge-Sort

- ◆ Merge-sort on an input sequence S with n elements consists of three steps:
 - **Divide**: If S has zero or one element, return S immediately. Otherwise, remove all the elements from S and put them into two sequences, S_1 and S_2 , each containing about half of the elements of S ; S_1 contains the first $\lceil n/2 \rceil$ elements of S and S_2 contains the remaining $\lfloor n/2 \rfloor$ elements
 - **Recur**: recursively sort S_1 and S_2
 - **Conquer**: merge S_1 and S_2 into a sorted sequence S
- ◆ The base case for the recursion are subproblems of size 0 or 1

$\lceil n/2 \rceil$ – ceiling of x – the smallest integer m , such that $x \leq m$, e.g.: $\lceil 5/2 \rceil = 3$
 $\lfloor n/2 \rfloor$ – floor of x – the largest integer k , such that $k \leq x$, e.g.: $\lfloor 5/2 \rfloor = 2$

Merge-sort

Algorithm *mergeSort*(S, C)

Input sequence S with n elements, comparator C

Output sequence S sorted
according to C

if $S.size() > 1$

$(S_1, S_2) \leftarrow \text{partition}(S, n/2)$

mergeSort(S_1, C)

mergeSort(S_2, C)

$S \leftarrow \text{merge}(S_1, S_2)$

Merging Two Sorted List-based Sequences

Algorithm *merge*(*A*, *B*, *S*)

Input sorted sequences *A* and *B* and an empty sequence *S* implemented as linked list

Output sorted sequence $S = A \cup B$

while $\neg A.isEmpty() \wedge \neg B.isEmpty()$

if *A.first().element()* \leq *B.first().element()*

 { move the first element of *A* at the end of *S* }

S.addLast(*A.remove*(*A.first()*))

else

 { move the first element of *B* at the end of *S* }

S.addLast(*B.remove*(*B.first()*))

 { move the remaining elements of *A* to *S* }

while $\neg A.isEmpty()$






S.addLast(*A.remove*(*A.first()*))

 { move the remaining elements of *B* to *S* }

while $\neg B.isEmpty()$

S.addLast(*B.remove*(*B.first()*))

return *S*



Merging Two Sorted Array-based Sequences

Algorithm *merge*(*A*, *B*, *S*)

Input sorted sequences *A* and *B* and an empty sequence *S*, all of which are implemented as arrays

Output sorted sequence $S = A \cup B$

$i \leftarrow j \leftarrow 0$

while $i < A.size() \wedge j < B.size()$ **do**

if $A.get(i) \leq B.get(j)$

 { copy *i*th element of *A* to the end of *S* }

S.addLast(*A.get*(*i*))

$i \leftarrow i+1$

else

S.addLast(*B.get*(*j*))

 { copy *j*th element of *B* to the end of *S* }

$j \leftarrow j+1$

 { move the remaining elements of *A* to *S* }

while $i < A.size()$ **do** *S.addLast*(*A.get*(*i*)); $i \leftarrow i+1$

 { move the remaining elements of *B* to *S* }

while $j < B.size()$ **do** *S.addLast*(*B.get*(*j*)); $j \leftarrow j+1$

return *S*

i) A

0	1	2
24	45	63

B

0	1
17	31

S

0	1	2	3	4

ii) A

0	1	2
24	45	63

B

0	1
17	31

S

0	1	2	3	4
17				

iii) A

0	1	2
24	45	63

B

0	1
17	31


S

0	1	2	3	4
17	24			

iv) ...


Merge-Sort Tree

- ◆ An execution of merge-sort is depicted by a binary tree
 - each node represents a recursive call of merge-sort and stores
 - ◆ unsorted sequence before the execution and its partition
 - ◆ sorted sequence at the end of the execution
 - the root is the initial call
 - the leaves are calls on subsequences of size 0 or 1
- ◆ Sequence: (7,2,9,4)




Execution Example (7,2,9,4,3,8,6,1)

◆ Partition




Recursion

◆ Recursive call, partition




Execution Example (cont.)

◆ Recursive call, partition




Execution Example (cont.)

◆ Recursive call, base case




Execution Example (cont.)

◆ Recursive call, base case



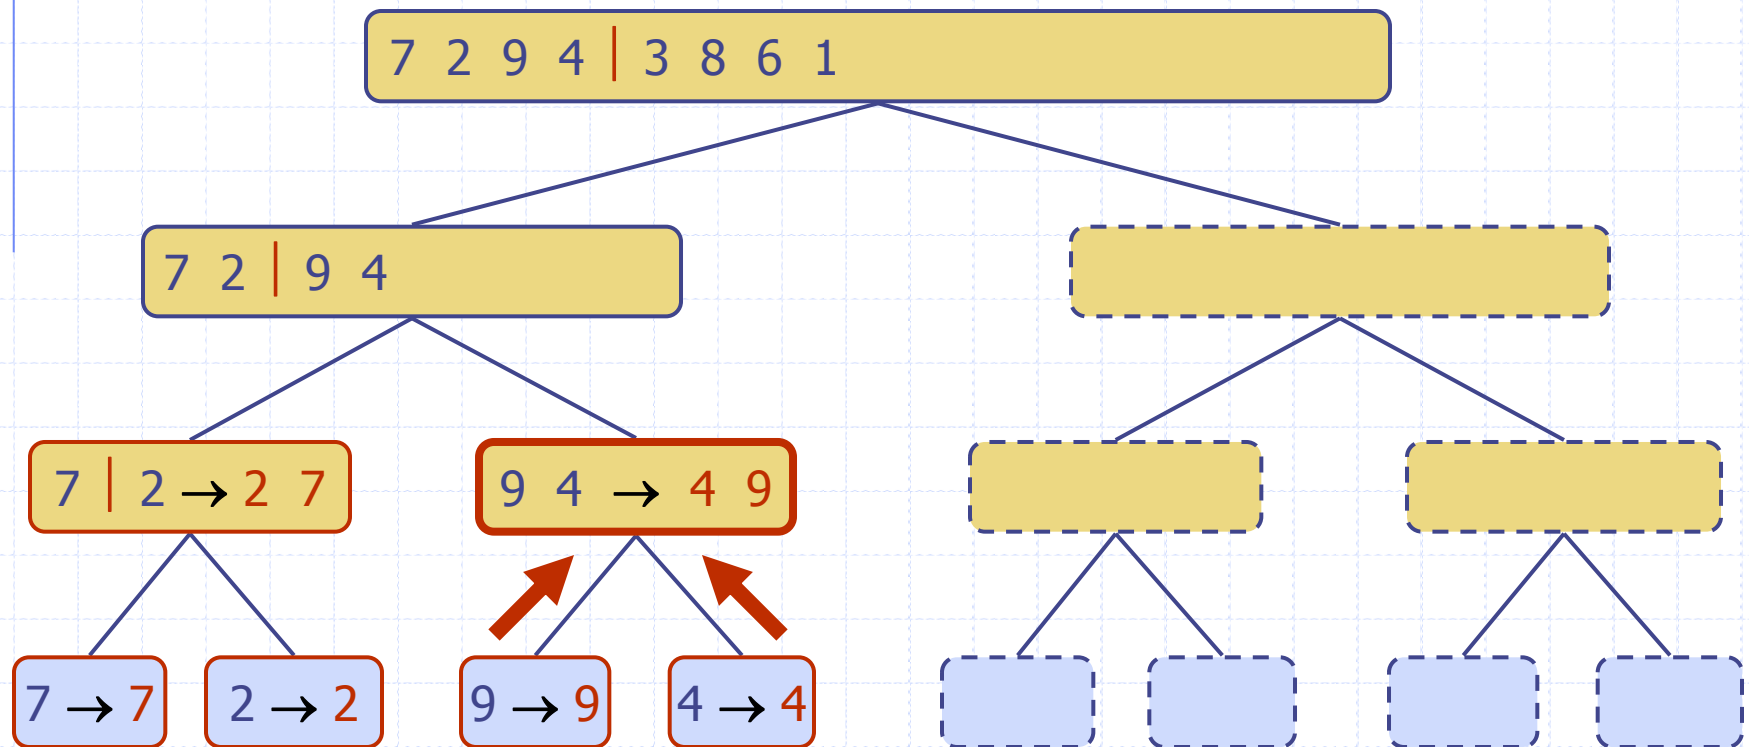
Execution Example (cont.)

◆ Merge



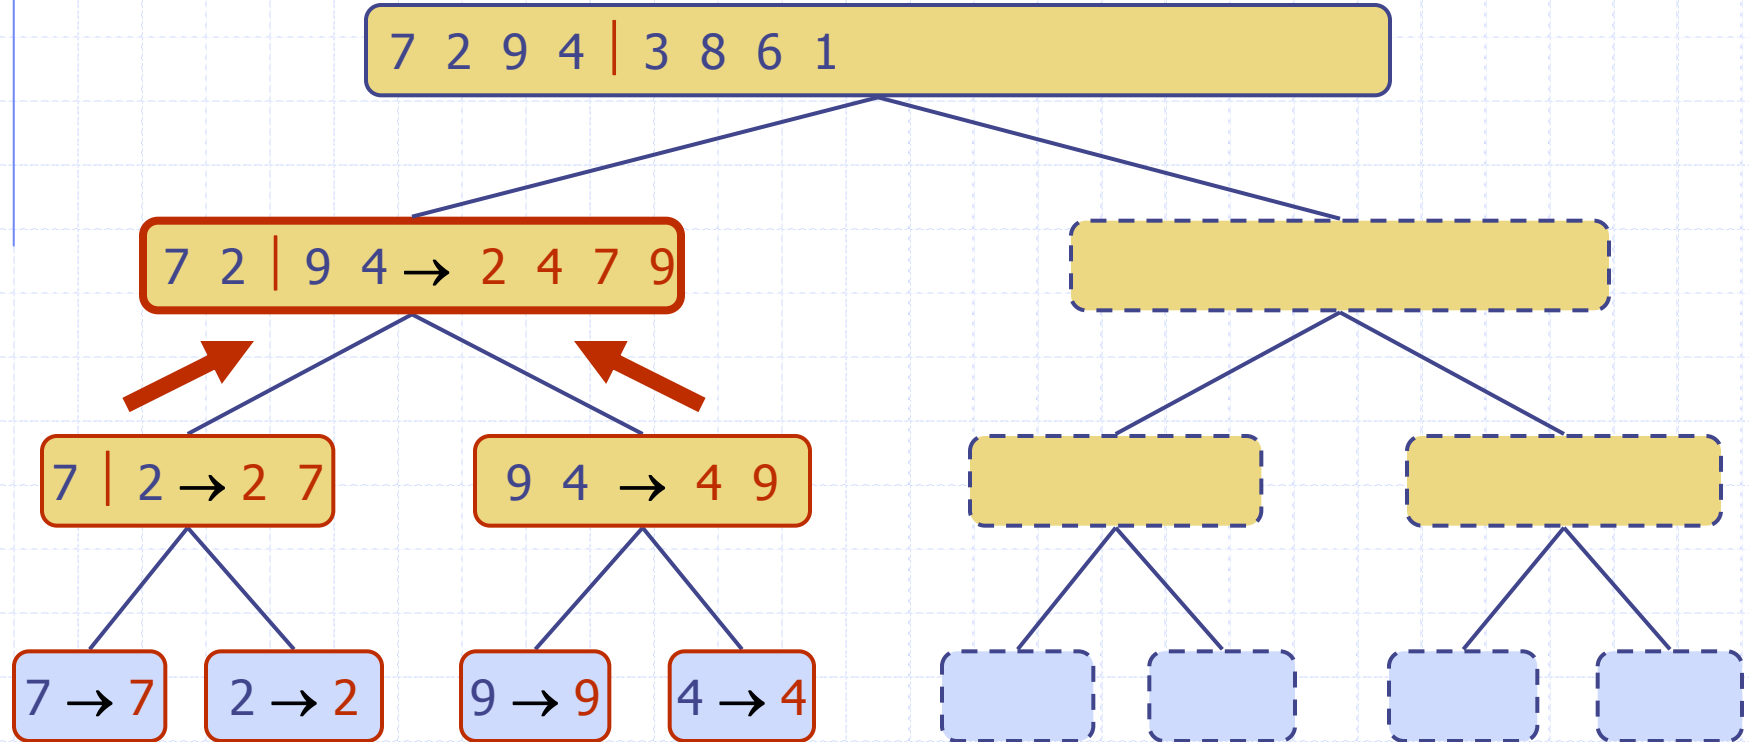
Execution Example (cont.)

◆ Recursive call, ..., base case, merge



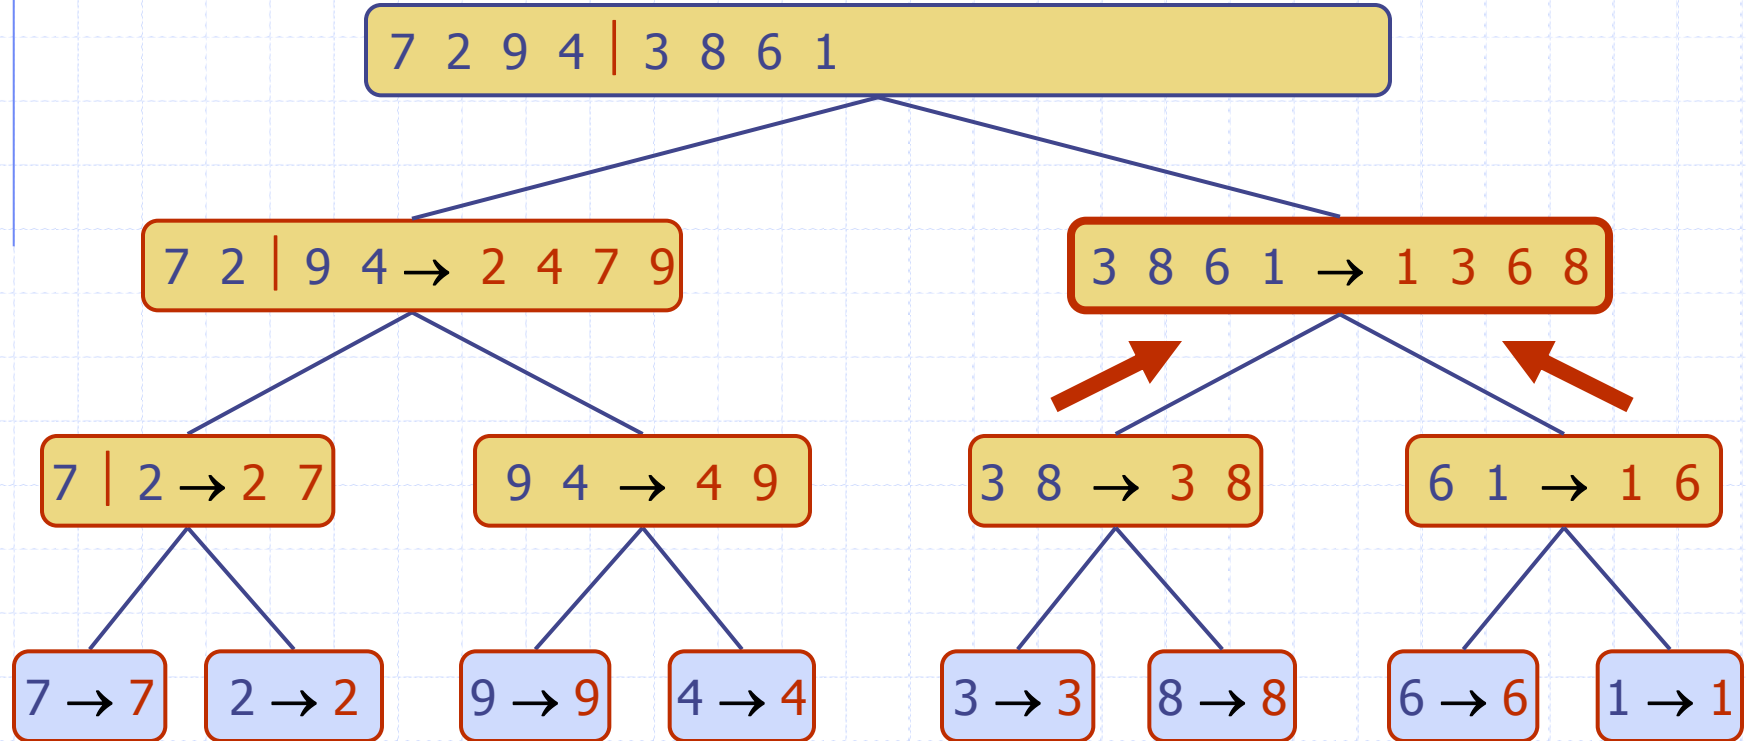
Execution Example (cont.)

◆ Merge




Execution Example (cont.)

◆ Recursive call, ..., merge, merge



Execution Example (cont.)

◆ Merge



Analysis of Merge-Sort

- ◆ The height h of the merge-sort tree is $O(\log n)$
 - at each recursive call we divide in half the sequence,
- ◆ The overall amount of work done at the nodes of depth i is $O(n)$
 - tree has 2^i nodes at depth i
 - we partition and merge 2^i sequences of size $n/2^i$
 - Overall time spent at all the nodes at depth i is $O(2^i \cdot n/2^i)$ which is $O(n)$
- ◆ Thus, the total running time of merge-sort is $O(n \log n)$


depth	#seqs	size
-------	-------	------

0	1	n
---	---	-----

1	2	$n/2$
---	---	-------


i	2^i	$n/2^i$
-----	-------	---------

...
-----	-----	-----



Merge Sort

Quick-Sort



Quick-Sort




◆ **Quick-sort** is a randomized sorting algorithm based on the divide-and-conquer paradigm:


- **Divide**: If S has at least two elements (otherwise nothing needs to be done) pick a random element x (called **pivot**). Remove all the elements from S and put into three sequences:

- ◆ L elements less than x
- ◆ E elements equal x
- ◆ G elements greater than x

- **Recur**: sort L and G
- **Conquer**: Put back the elements into S in order by first inserting the elements of L , then those of E and finally those of G .

◆ **Common Practice**: choose the pivot to be the last element in S





Partition

- ◆ We partition an input sequence as follows:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the comparison with the pivot x
- ◆ Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$ time
- ◆ Thus, the partition step of quick-sort takes $O(n)$ time

Algorithm *partition*(S, p)

Input sequence S , position p of pivot
Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

$L, E, G \leftarrow$ empty sequences

$x \leftarrow S.remove(p)$

$E.addLast(x)$

while $\neg S.isEmpty()$

$y \leftarrow S.remove(S.first())$

if $y < x$

$L.addLast(y)$

else if $y = x$

$E.addLast(y)$


else $\{ y > x \}$

$G.addLast(y)$

return L, E, G


Quick-Sort Tree

- ◆ An execution of quick-sort is depicted by a binary tree
 - Each node represents a recursive call of quick-sort and stores
 - ◆ Unsorted sequence before the execution and its pivot
 - ◆ Sorted sequence at the end of the execution
 - The root is the initial call
 - The leaves are calls on subsequences of size 0 or 1




Execution Example (7,2,9,4,3,7,6,1)

◆ Pivot selection



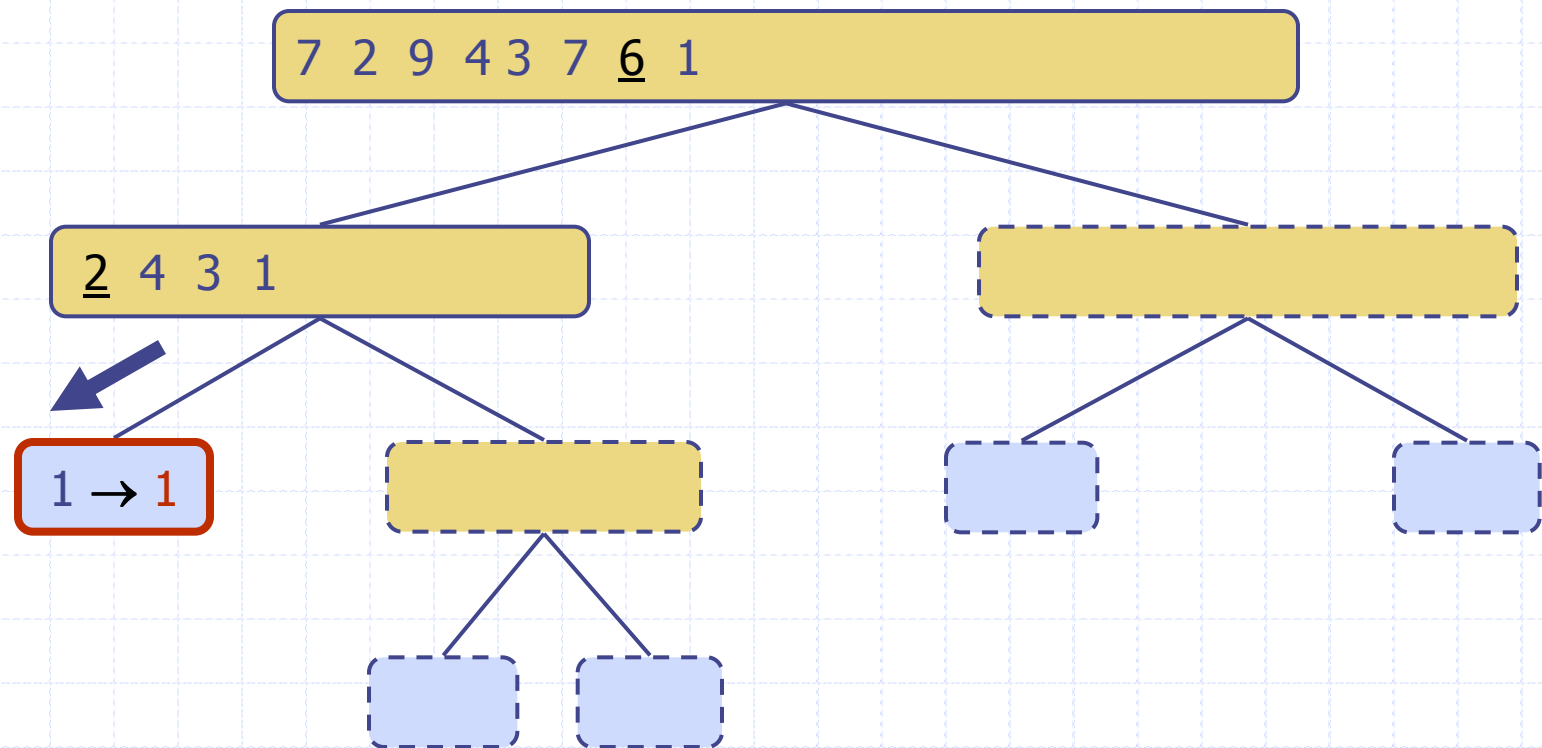
Execution Example (cont.)

◆ Partition, recursive call, pivot selection



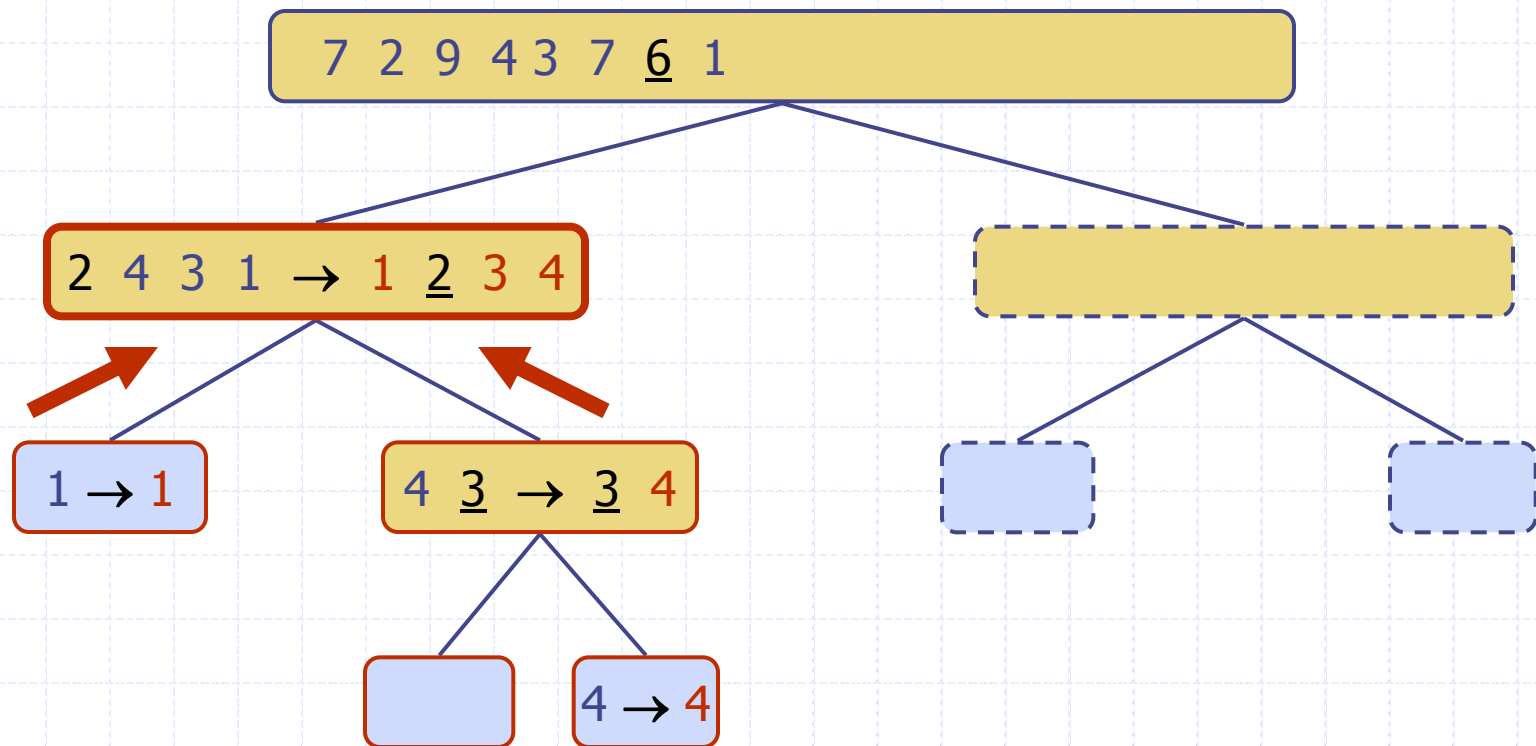
Execution Example (cont.)

◆ Partition, recursive call, base case



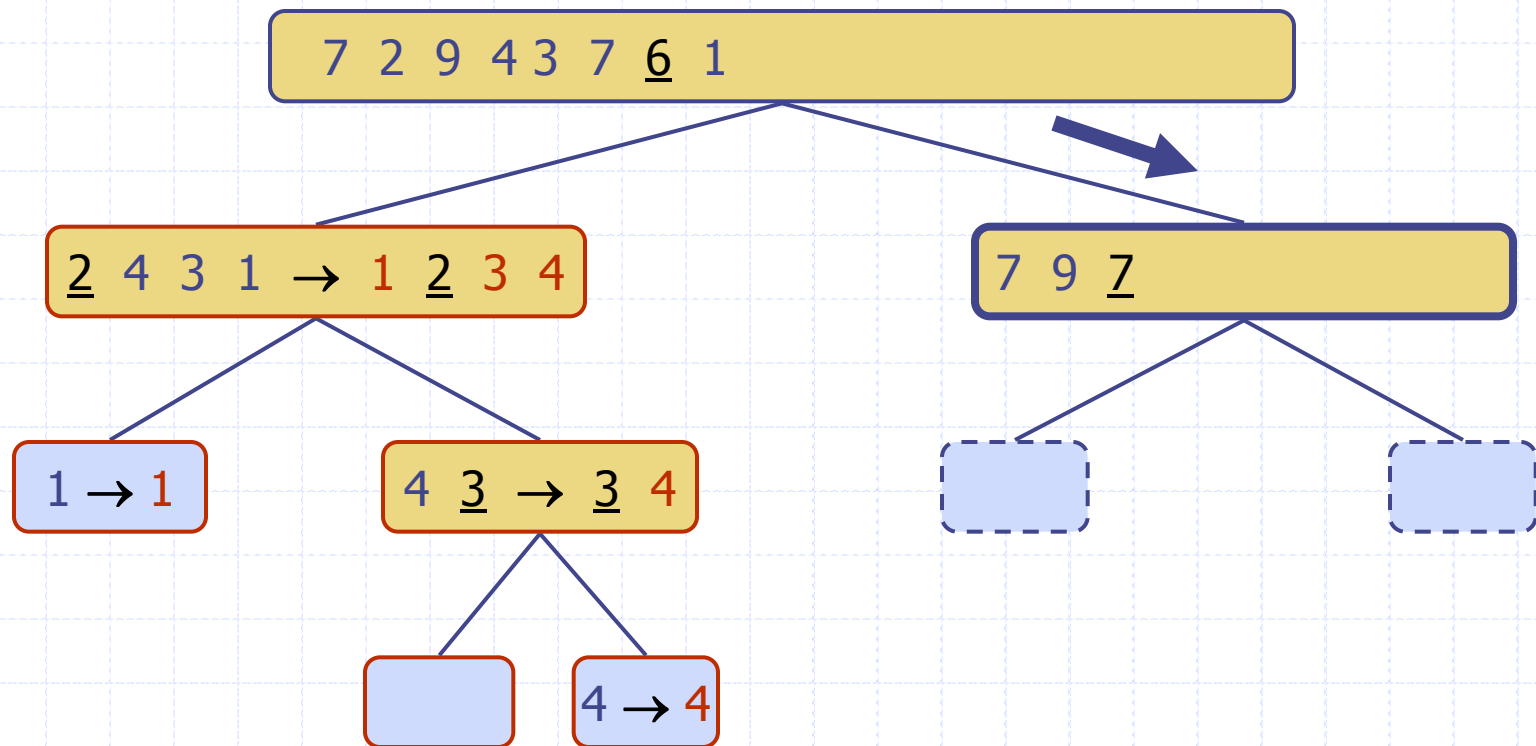
Execution Example (cont.)

◆ Recursive call, ..., base case, join



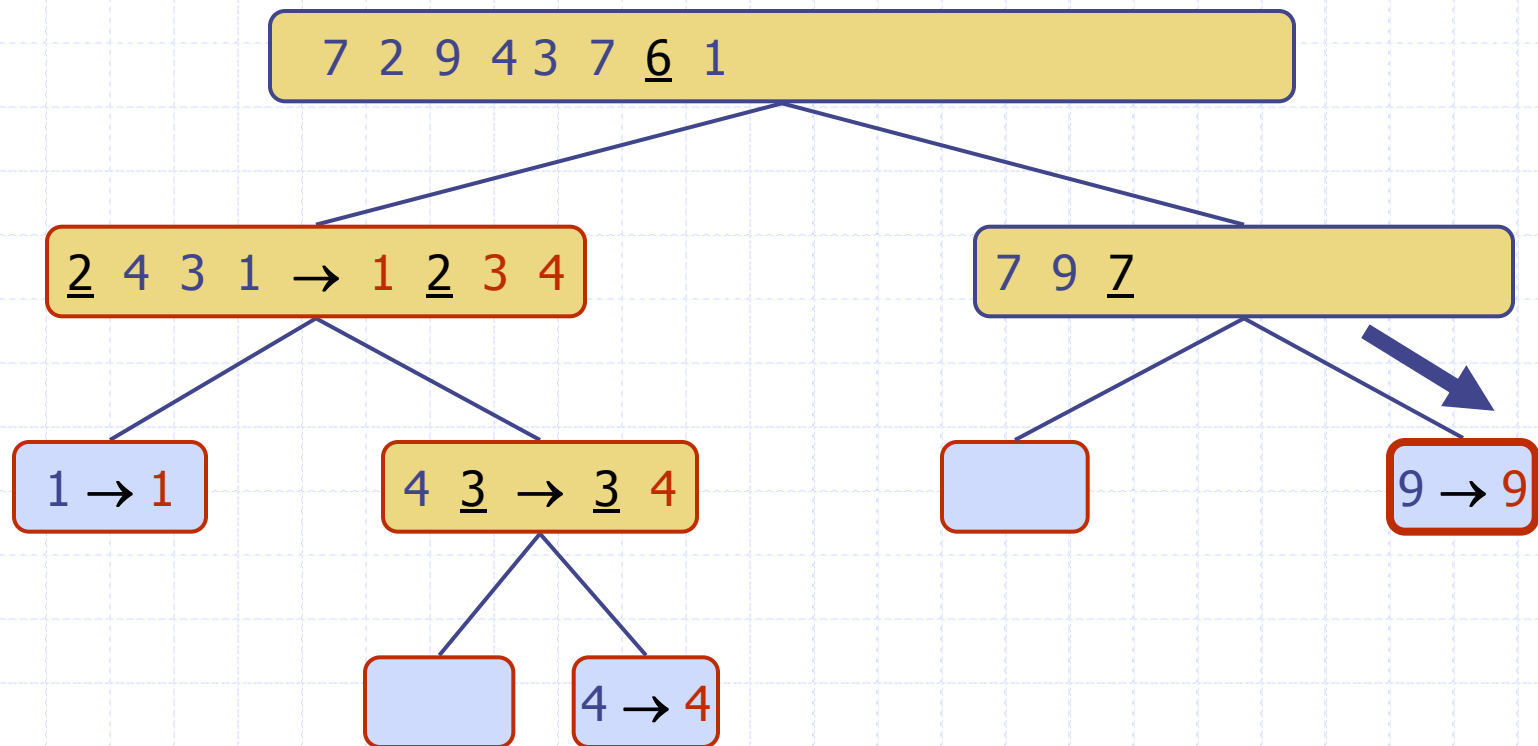
Execution Example (cont.)

◆ Recursive call, pivot selection



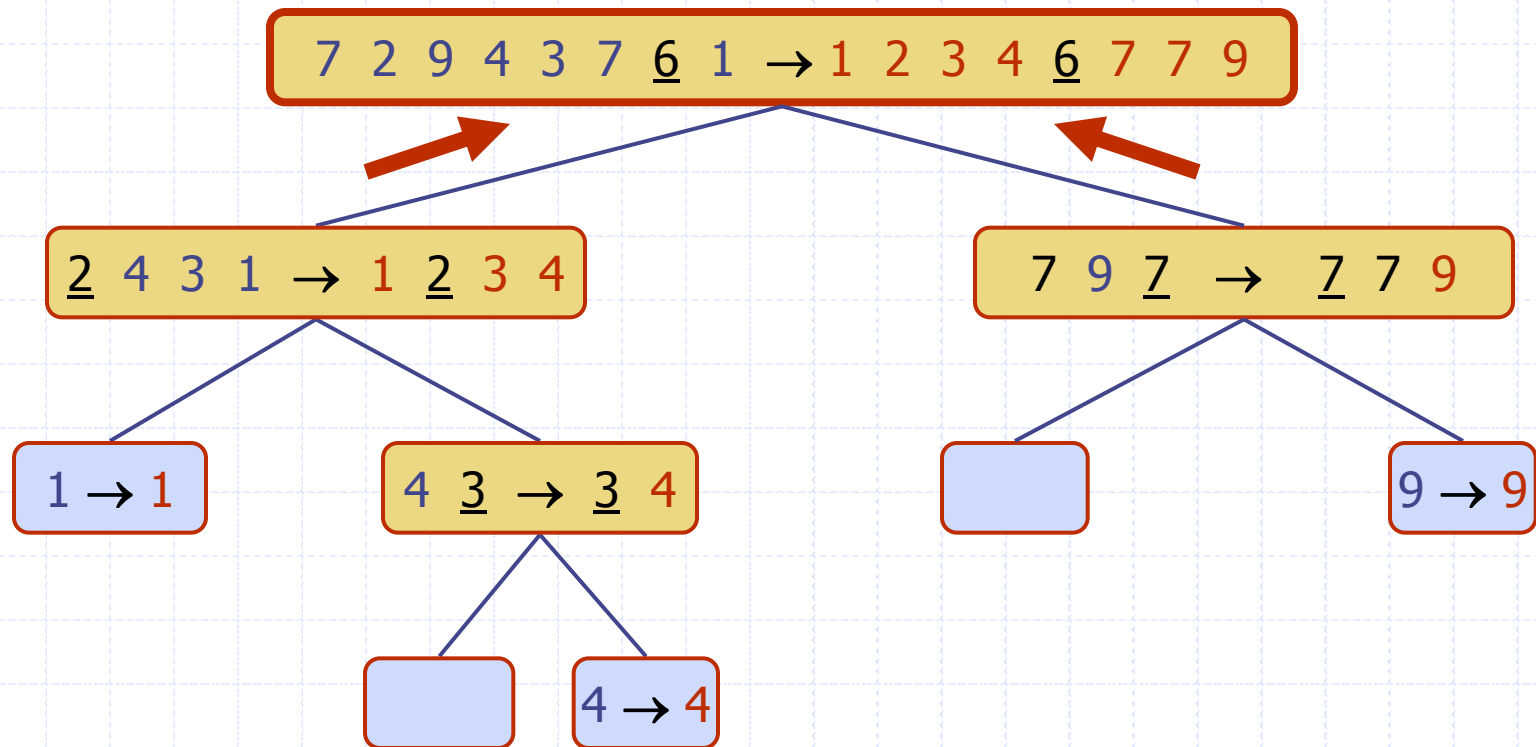
Execution Example (cont.)

◆ Partition, ..., recursive call, base case



Execution Example (cont.)

◆ Join, join



Worst-case Running Time

- ◆ The worst case for quick-sort occurs when the pivot is the unique minimum or maximum element
- ◆ One of L and G has size $n - 1$ and the other has size 0
- ◆ The running time is proportional to the sum
$$n + (n - 1) + \dots + 2 + 1$$
- ◆ Thus, the worst-case running time of quick-sort is $O(n^2)$

depth time

0

n

1


$n - 1$

...

...


$n - 1$

1




Expected Running Time

- ◆ Consider a recursive call of quick-sort on a sequence of size s
 - **Good call:** the sizes of L and G are each less than $3s/4$
 - **Bad call:** one of L and G has size greater than $3s/4$




Good call




Bad call

- ◆ A call is **good** with probability $1/2$
 - $1/2$ of the possible pivots cause good calls:



Expected Running Time, Part 2

- ◆ **Probabilistic Fact:** The expected number of coin tosses required in order to get k heads is $2k$
- ◆ For a node of depth i , we expect
 - $i/2$ ancestors are good calls
 - The size of the input sequence for the current call is at most $(3/4)^{i/2}n$
- ◆ Therefore, we have
 - For a node of depth $2\log_{4/3}n$, the expected input size is one
 - The expected height of the quick-sort tree is $O(\log n)$
- ◆ The amount of work done at the nodes of the same depth is $O(n)$
- ◆ Thus, the expected running time of quick-sort is $O(n \log n)$



Exercise – Merge-sort and Quick-sort


◆ Perform merge-sort and quick-sort on the following sequence of numbers:

(3, 5, 1, 9, 3)


In-Place Quick-Sort – optimisation of “divide” step



- ◆ Quick-sort can be implemented to run in-place
- ◆ “Divide” step can be done “in place”, that is, without using any additional array, in the following way.
- ◆ Assume we want to divide $A[\text{leftEnd} \dots \text{rightEnd}]$ with respect to pivot $A[\text{rightEnd}]$.
- ◆ Maintain two indices, L (left cursor) and R (Right cursor), with initial values leftEnd and $\text{rightEnd}-1$, respectively. The elements which have not been considered yet are in $A[(L+1) \dots (R-1)]$.



In-Place Quick-Sort – optimisation of “divide” step




□ Iterate until $L \leq R$:

- Keep increasing L by 1 until an element $A[L]$ is smaller than the pivot and $L \leq R$.
- Keep decreasing R by 1 until an element $A[R]$ is larger than the pivot and $L \leq R$.
- If $L < R$ swap $A[L]$ and $A[R]$, and proceed to the next iteration.

□ Swap $A[L]$ and pivot


Iteration 1:

$L < R$ – swap $A[L]$ and $A[R]$




Iteration 2:

$L < R$ – swap $A[L]$ and $A[R]$




Iteration 3:

$L < R$ – swap $A[L]$ and $A[R]$




$L > R$ – Swap $A[L]$ and pivot




Summary of Sorting Algorithms

Algorithm	Time	Notes
selection-sort	$O(n^2)$	<ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs)
insertion-sort	$O(n^2)$	<ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs)
quick-sort	$O(n \log n)$ expected	<ul style="list-style-type: none">▪ in-place▪ fastest (good for large inputs)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ in-place▪ fast (good for large inputs)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ sequential data access▪ fast (good for huge inputs)

Bucket-Sort





Bucket-Sort

- ◆ Bucket-sort does not use comparison
- ◆ Let S be a sequence of n (key, element) entries with keys in the range $[0, N - 1]$
- ◆ Bucket-sort uses the keys as indices into an auxiliary array B of sequences (buckets)

Phase 1: Empty sequence S by moving each entry (k, o) into its bucket $B[k]$

Phase 2: For $i = 0, \dots, N - 1$, move the entries of bucket $B[i]$ to the end of sequence S

- ◆ Analysis:
 - Phase 1 takes $O(n)$ time
 - Phase 2 takes $O(n + N)$ time
- Bucket-sort takes $O(n + N)$ time

Algorithm *bucketSort*(S, N)

Input sequence S of (key, element) items with keys in the range $[0, N - 1]$

Output sequence S sorted by increasing keys

$B \leftarrow$ array of N empty sequences

while $\neg S.isEmpty()$

$f \leftarrow S.first()$

$(k, o) \leftarrow S.remove(f)$

$B[k].addLast((k, o))$

for $i \leftarrow 0$ **to** $N - 1$


while $\neg B[i].isEmpty()$

$f \leftarrow B[i].first()$

$(k, o) \leftarrow B[i].remove(f)$

$S.addLast((k, o))$


Example



◆ Key range [0, 9]



Phase 1



Phase 2



Exercise – Bucket-sort


◆ Perform bucket-sort on the following sequence of numbers:

(3, 5, 1, 9, 3, 7, 8, 8)

Stability of sorting

- ◆ Sorting algorithm is stable if the relative order of any two items with the same key in an input sequence is preserved after the execution of the algorithm.
- ◆ Stable sorting algorithms:
 - Merge-sort
 - Bucket-sort
- ◆ Unstable sorting algorithms:
 - Quick-sort


Sorting Lower Bound



Comparison-Based Sorting




- ◆ Many sorting algorithms are comparison based.
 - They sort by making comparisons between pairs of objects
 - Examples: selection-sort, insertion-sort, heap-sort, merge-sort, quick-sort
- ◆ Let us therefore derive a lower bound on the running time of any algorithm that uses comparisons to sort n elements, x_1, x_2, \dots, x_n .




Counting Comparisons


- ◆ Let us just count comparisons then.
- ◆ Each possible run of the algorithm corresponds to a root-to-leaf path in a **decision tree**



Decision Tree Height

- ◆ The height of the decision tree is a lower bound on the running time
- ◆ Every input permutation must lead to a separate leaf output
- ◆ If not, some input ...4...5... would have same output ordering as ...5...4..., which would be wrong
- ◆ Since there are $n! = 1 \cdot 2 \cdot \dots \cdot n$ leaves, the height is at least $\log(n!)$





The Lower Bound

- ◆ Any comparison-based sorting algorithm takes at least $\log(n!)$ time
- ◆ Therefore, any such algorithm takes time at least

$$\log(n!) \geq \log \left(\frac{n}{2} \right)^{\frac{n}{2}} = (n/2) \log(n/2).$$

- ◆ That is, any comparison-based sorting algorithm must run in $\Omega(n \log n)$ time.