

Security on the Internet



Outline

- ▶ Internet security threats
- ▶ HTTP authentication
- ▶ Hash functions and encryption
- ▶ Public key encryption
- ▶ Pretty Good Privacy
- ▶ Transport Layer Security and HTTPS
- ▶ Digital certificates



Pervasive Computing

- ▶ social media
- ▶ phones
- ▶ cars
- ▶ medical devices
- ▶ cash

- ▶ The Internet of Things (IoT)



Computer Security

- ▶ The engineering of systems that exhibit the following properties:
 - ▶ Confidentiality
 - ▶ Integrity
 - ▶ Availability



Confidentiality

- ▶ Non-disclosure of information except to another authorised agent



Integrity

Every piece of data is as the last authorised modifier left it.

- ▶ Data integrity:

ensuring that the data has not been deleted or altered without permission.

- ▶ Software integrity:

ensuring that the software programs have not been altered, whether by an error, a malicious user, or a virus.



Availability

- ▶ being accessible and usable upon demand by an authorized entity.



Computer Security and the Internet

The internet can make attacks easier because of:

- ▶ Action at a Distance
- ▶ Technique Propagation
- ▶ Automation, e.g. through harnessing distributed computation



Vulnerabilities and exploits

- ▶ Most software systems do **not** exhibit these properties.
- ▶ When a software system is not secure, it has one or more *vulnerabilities*.
- ▶ Malicious users attack vulnerable systems through *exploits*.
- ▶ There is a one-to-many mapping between vulnerabilities and exploits.
- ▶ An exploit is a piece of code, or a replicable procedure, which is able to exploit the vulnerability.
- ▶ A given vulnerability may have zero or more known exploits.



Vulnerability Announcements

- ▶ All software has vulnerabilities.
- ▶ Some of these vulnerabilities are known, others are not.
- ▶ Of the known vulnerabilities, different parties may have different knowledge.
- ▶ When a vulnerability is detected by the security community, e.g. CERT, software vendors are informed.
- ▶ There is a typically grace period before the vulnerability is announced to the general public.
 - ▶ <https://www.us-cert.gov/ncas/alerts.xml>
 - ▶ <https://web.nvd.nist.gov/view/vuln/search>



Zero-Day Vulnerabilities

- ▶ Vulnerabilities that have not been publicly unannounced are called "zero-day".
- ▶ Attackers users who have knowledge of a zero-day vulnerability and a corresponding exploit are in a very powerful position.
- ▶ Government agencies bid for zero-day vulnerabilities and exploits on the black market



Authentication

- ▶ In order to defend our system from unauthorised users, we must distinguish between authorised and unauthorised users.
- ▶ The process of determining the true identity of a user is called authentication.
- ▶ The simplest technique for authentication is to use a secret password or pass-phrase.



Access control

- ▶ Preventing illegitimate access can be split into two issues
 - ▶ **Authentication:** Determining who is trying to gain access to your host
 - ▶ **Access Control:** Determining whether that individual is allowed to access a resource on your host



Web Authentication

- ▶ Authentication mechanisms of web servers prevent illegitimate access to resources
- ▶ Resources are often grouped into named **realms**, which users can be allowed access to
- ▶ They require clients to demonstrate who they are, from extra data sent with their messages



Authentication HTTP status

- ▶ If a client tries to access a secured web server with no authentication, it receives a response with status code:

401 Authentication Required

- ▶ This tells the client to supply identification
- ▶ The response will contain a field:
WWW-Authenticate: ...
- ▶ specifying the authentication scheme required
- ▶ The client provides authentication in new request



Proof of identity

- ▶ Common forms of proof of identity are:
 - ▶ Username plus password
 - ▶ Public-Key Cryptography and Digital certificates
- ▶ Software applications (agents) must also authenticate themselves, and may have identities different from but based on that of their owner



HTTP basic authentication

- ▶ Basic authentication is indicated by the Basic scheme being passed in the 401 response:

WWW-Authenticate: Basic realm="somerealm"

- ▶ Basic secure request gives credentials as base64 encoding of text

username:password

- ▶ Encoding put into Authorization request field:

Authorization: Basic QWxhZGR=



Hash functions

- ▶ A **hash** of some data is a transformation of that data into a fixed length string, from which the original cannot be deduced:
 - ▶ I.e. it is a one-way function
- ▶ Use a hash wherever data across places or times need to be compared to check they are the same: if hash values match this is **almost** proof
- ▶ Value of algorithm depends on:
 - ▶ Difficulty of deducing original from digest
 - ▶ Lack of **collisions**: 2 messages with same digests



Password Hashing

- ▶ When a server authenticates a user using a password, it compares the supplied password against a field stored in the database.
- ▶ If we store the original password, then if an attacker breaks into the server, they can steal all the users' passwords.
- ▶ Passwords are often reused across different servers.
- ▶ To prevent this, we store a hash of the password instead.
- ▶ We also include random salt (also called a nonce) to prevent rainbow attacks.
- ▶ On Unix, the file `/etc/shadow` contains hashed passwords.



Password authentication

- ▶ The client sends the username and password (in some form) with every request, and the web server matches against its list to see whether access is allowed.
- ▶ Eavesdropping is a problem with this if the authentication occurs over an unencrypted channel.



MD5

- ▶ MD5 is a hash function devised by Ronald L. Rivest
- ▶ MD stands for Message Digest (another word for hash).
- ▶ It takes an arbitrarily long string and produces a string of fixed length, the hash value or *digest*.
- ▶ Collisions are possible in MD5.
- ▶ Can try it out with UNIX command md5sum.

"digest" abfd2c0ecb4e9dec4a6b1159d5fea334

"digest" 5a20c77381e982467465dd18facf0807

"digest " e21681785dc42cfc30867e4fcf78edaf



SHA-2

- ▶ MD5 should *no longer used* for cryptographic applications, as it was shown to be insecure.
- ▶ The current standard is SHA 2
- ▶ MD5 is still considered to be secure for password hashing in the case of legacy applications
 - ▶ New applications should make use of SHA-2



HTTP digest authentication

1. Client tries to access a realm on the server
2. Server responds asking for authentication (HTTP 401) using digests and providing a unique identifier, a number used once (nonce), for the request
3. Client sends a digest of the concatenation of:
 - ▶ username, realm, password, URL, request method and nonce id



HTTP digest authentication

- ▶ Digest authentication is indicated by the Digest scheme passed in the 401 response:

**WWW-Authenticate: Digest
realm="somerealm" algorithm="MD5"
nonce="564dsd" ...**

- ▶ Create digest of password plus other data
- ▶ Digest put into Authorization request field:

**Authorization: Digest Username="Simon",
response="6629fae49393a053",
realm="somerealm", nonce="564dsd" ...**



Eavesdropping

- ▶ Eavesdropping is a problem on the internet:
 - ▶ Happens at any point in route between hosts
 - ▶ Whole physical network cannot (realistically) be secured
- ▶ Promiscuous Mode: Ethernet allows hosts to request that all data sent through a network be passed to it, even if not a router or destination.
- ▶ A sniffer is a device/program to monitor all data across a network, and so can be used to extract passwords etc.
- ▶ **Encryption** is the transformation of data to a form unreadable by anyone but the intended recipient.
- ▶ The algorithm for performing the transformation is called a **cryptographic cipher**.



Encryption types

- ▶ Encryption on the internet takes several forms:
 - ▶ **Link encryption:** Encrypts all communication across a physical link, but it is expensive and unrealistic over large scale
 - ▶ **Document encryption:** Documents encrypted, sent, then decrypted by document-handling applications
 - ▶ **Transport Layer Security:** TLS encrypts all messages at the TCP layer



Hashes and encryption

- ▶ Both a hash and an encryption of a message are a transformation of that message into some new data that gives no clue to the original.
- ▶ The original message can be computed from the encrypted data (decrypted) given the right information, but this is not true of a hash.



Ciphers

- ▶ A cryptographic cipher provides an encryption

$$C = E(M, K)$$

- ▶ And a decryption function:

$$M = D(C, K)$$

- ▶ The original message is called the plain-text.
- ▶ The encrypted message is called the cipher-text.
- ▶ The decryption function can only be computed if you have access to the correct secret key K.



Single-key encryption

- ▶ A **key** is a secret piece of data that sender and receiver use to encrypt and decrypt messages
- ▶ As no-one else knows the key, no-one else can decrypt the message
- ▶ In single key encryption, the same key is used by both parties to both encrypt and decrypt.
- ▶ Another name for single-key encryption is symmetric cryptography.
- ▶ An example of a symmetric cipher is the Advanced Encryption Standard (AES).



Key Sizes

- ▶ The space of possible keys must be large.
- ▶ This is in order to prevent an attacker simply trying every possible key
 - ▶ This is called a brute-force attack.
- ▶ For a symmetric cipher such as AES:
 - ▶ Maximum key size is 256 bits,
 $2^{256} =$
11579208923731619542357098500868790785326998466564056
4039457584007913129639936 possible keys.



Key Generation

- ▶ Keys must be hard to guess.
- ▶ Therefore they should be randomly chosen.
- ▶ It is very hard to generate genuinely random data using a deterministic digital computer.
- ▶ Usually the best we can do is to use a Pseudo-Random Number Generator (PRNG).
- ▶ In cryptography we make use of external sources of randomness (entropy) such as heat in order to create Secure-Random Number Generators.
- ▶ On Unix, see `/dev/urandom`



Key Management

Problems with symmetric ciphers:

1. We need a separate key for every possible pair of users.
 2. We need to agree on a key with the other party.
- ▶ We cannot simply send the key, because of eavesdroppers.

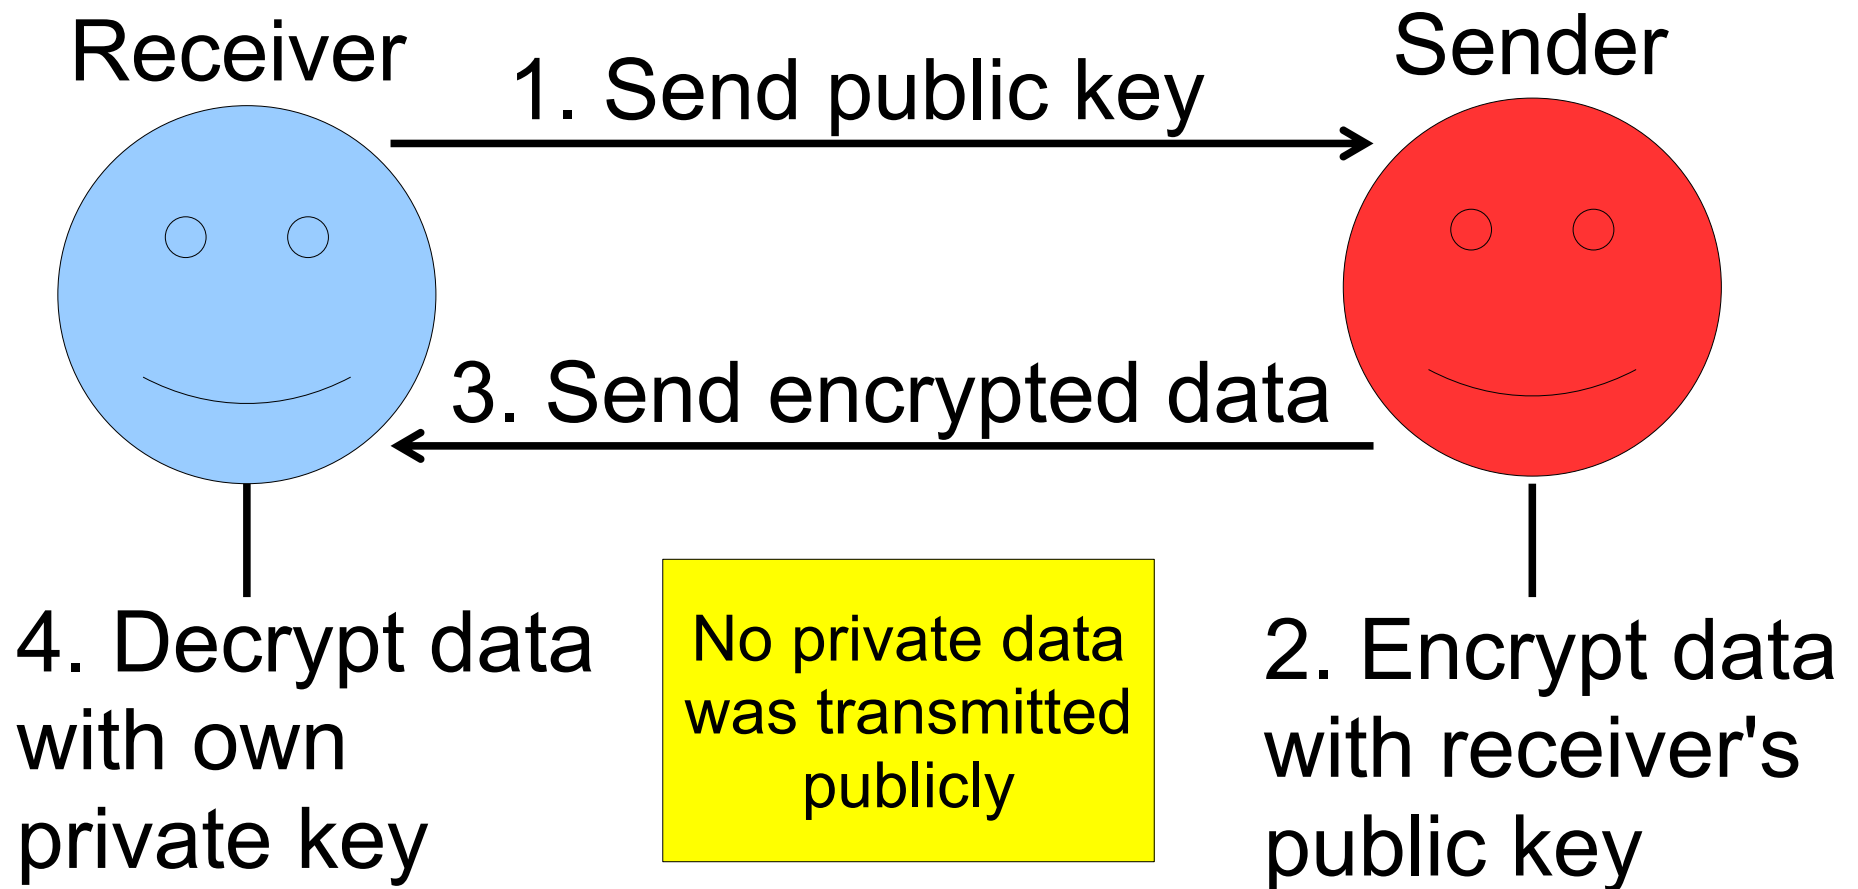


Public key encryption

- ▶ In public key cryptography, each user has a pair of keys: a **public** one and a **private** one.
- ▶ These are called *asymmetric* ciphers.
- ▶ The public key can be made available to anyone wanting to send an encrypted message to the user
- ▶ Sender encrypts their messages with the public key: the algorithm ensures the message can only be decrypted with the private key
- ▶ On receiving an encrypted message, the user uses the private key to decrypt the messages



Scenario 1: Secret message



RSA

- ▶ One of the best known public key encryption algorithms is RSA, named after the inventors: Rivest, Shamir and Adleman
- ▶ The idea of RSA is to use two very large prime numbers for the keys



RSA key generation

1. Generate two large primes: p, q
2. Calculate the product: $n = pq$
3. Calculate the totient: $m = (p - 1)(q - 1)$
4. Find a co-prime to totient m : e
5. Choose integers d, i so that: $d = (1 + im) / e$
6. The public key is (n, e)
7. The private key is (n, d)



RSA generation example

1. Generate two large primes: $7, 11$ (p, q)



RSA generation example

- | | | |
|-------------------------------|---------|------------|
| 1. Generate two large primes: | $7, 11$ | (p, q) |
| 2. Calculate the product: | 77 | $(n = pq)$ |



RSA generation example

- | | | |
|-------------------------------|--------------|----------------------|
| 1. Generate two large primes: | <i>7, 11</i> | (p, q) |
| 2. Calculate the product: | <i>77</i> | $(n = pq)$ |
| 3. Calculate the totient: | <i>60</i> | $m = (p - 1)(q - 1)$ |



RSA generation example

- | | | |
|--------------------------------|-------|--------------------|
| 1. Generate two large primes: | 7, 11 | (p, q) |
| 2. Calculate the product: | 77 | (n = pq) |
| 3. Calculate the totient: | 60 | m = (p - 1)(q - 1) |
| 4. Find a co-prime to totient: | 7 | e |

To find co-prime, count up each prime, and find first not divisible into totient:

Is 60 divisible by 2? Yes

Is 60 divisible by 3? Yes

Is 60 divisible by 5? Yes

Is 60 divisible by 7? No



RSA generation example

- | | | |
|--------------------------------|--------------------------------|--------------------|
| 1. Generate two large primes: | 7, 11 | (p, q) |
| 2. Calculate the product: | 77 | (n = pq) |
| 3. Calculate the totient: | 60 | m = (p - 1)(q - 1) |
| 4. Find a co-prime to totient: | 7 | e |
| 5. Choose integers i, d: | 5, 43 | |
| | $(1 + (5 \times 60)) / 7 = 43$ | $d = (1 + im) / e$ |

To find i and d, count up i from 1 until find first
(1 + (i x 60)) divisible by 7:

(1 + (1 x 60)) divisible by 7? No

(1 + (2 x 60)) divisible by 7? No

...

(1 + (5 x 60)) divisible by 7? Yes, d = 43



RSA generation example

- | | | |
|----------------------------------|--------------------------------|----------------------|
| 1. Generate two large primes: | 7, 11 | (p, q) |
| 2. Calculate the product: | 77 | $(n = pq)$ |
| 3. Calculate the totient: | 60 | $m = (p - 1)(q - 1)$ |
| 4. Find a co-prime to totient: | 7 | e |
| 5. Choose integers i, d : | 5, 43 | |
| | $(1 + (5 \times 60)) / 7 = 43$ | $d = (1 + im) / e$ |
| 6. The public key is (n, e) : | (77, 7) | (n, e) |
| 7. The private key is (n, d) : | (77, 43) | (n, d) |



RSA encryption algorithm

1. Obtain receiver's public key (n, e)
2. Convert message to an array of bits representing a large integer
 $M < 2^n$
3. Compute encrypted values $C = M^e \bmod n$

-
1. Obtain receiver's public key $(77, 7)$
 2. Represent data as +ve integer 6
 3. Compute encrypted value $6^7 \bmod 77 = 41$



RSA decryption algorithm

- | | |
|------------------------------|-------------------|
| 1. Use private key | (n, d) |
| 2. Receive encrypted message | C |
| 3. Calculate original value | $M = C^d \bmod n$ |

-
- | | |
|------------------------------|------------------------|
| 1. Use private key | $(77, 43)$ |
| 2. Receive encrypted message | 41 |
| 3. Calculate original value | $41^{43} \bmod 77 = 6$ |



RSA signing

- ▶ RSA can also be used to digitally sign a document.
- ▶ The signer uses their private key (n, d) to produce a signature.
- ▶ The verifier can check this using the signer's public key.
- ▶ The signature S of a message M is:

$$S = M^d \bmod n$$

- ▶ To verify the signature, the verifier checks that:

$$S^e = M \bmod n$$

- ▶ We typically hash the message, and then sign the hash.
 - ▶ MD5 is not suitable for digital signatures- use SHA-2 or SHA-3.
-



Hybrid crypto-systems

- ▶ For large messages RSA is typically used in combination with a symmetric cipher such as AES.
- ▶ RSA can be used to encrypt messages containing the symmetric keys, which can then be distributed over a public channel.
- ▶ The symmetric key is typically a temporary key that is only valid for a particular session.



Transport Layer Security

- ▶ Secure Socket Layer (SSL) was developed by Netscape Communications
- ▶ It operates between host-to-host protocols (TCP) and the application layer protocols (e.g. HTTP)
- ▶ For each communication, SSL uses the most recent secure communication protocol that both hosts can support
- ▶ Transport Layer Security (TLS) is a more recent variation on SSL, standardised by IETF



TLS negotiating protocol

- ▶ TLS initiates a cryptographic protocol between hosts with a Hello message
- ▶ Both parties declare what they can support and the strongest encryption available is chosen
- ▶ This allows for the change and development of encryption methods
- ▶ The client will choose the stronger of the two protocols




TLS sharing certificates

- ▶ The hosts then exchange **certificates**
- ▶ Digital certificates provide verifiable host data for authentication
- ▶ They also provide public keys for encrypting the communication
- ▶ We'll say more on certificates later
- ▶ The public keys in the certificates are used to encrypt communication over TLS



Man-in-the-middle attacks

- ▶ TLS tackles the problem of man-in-the-middle attacks
- ▶ A malicious host routes communication through itself, without being apparent to either sender or receiver
- ▶ Data can then be copied and, if not encrypted, read



TLS and the internet architecture

- ▶ TLS operates over TCP, but under HTTP or other application protocols
- ▶ Inserts a new layer into the four-layer internet layering model
- ▶ This layer deals with the issue of secure communication of application data



HTTPS

- ▶ HTTPS is HTTP over SSL/TLS
- ▶ Uses its own URI scheme
- ▶ https:...
- ▶ Has a different default TCP port (443)
- ▶ Otherwise the same as HTTP over TCP
- ▶ An HTTPS web server must have a digital certificate that it can use to authenticate itself with a client



Digital certificates

- ▶ A digital certificate is a block of data about a communicating host that is signed
- ▶ Signing a certificate means adding an encrypted hash of the host data, so that other hosts can check that:
 - ▶ You are who you say you are
 - ▶ The host data has not been tampered with



Host data

- ▶ **Host data can include:**
 - ▶ Public key
 - ▶ Validity period of certificate
 - ▶ URL of revocation centre
 - ▶ Name, institution, email address of owning user
 - ▶ Public key is used to secure communication to the host
 - ▶ Certificates are revoked if they are suspected of being compromised (like credit cards)
 - ▶ Revocation centres provide lists of revoked certificates to check against

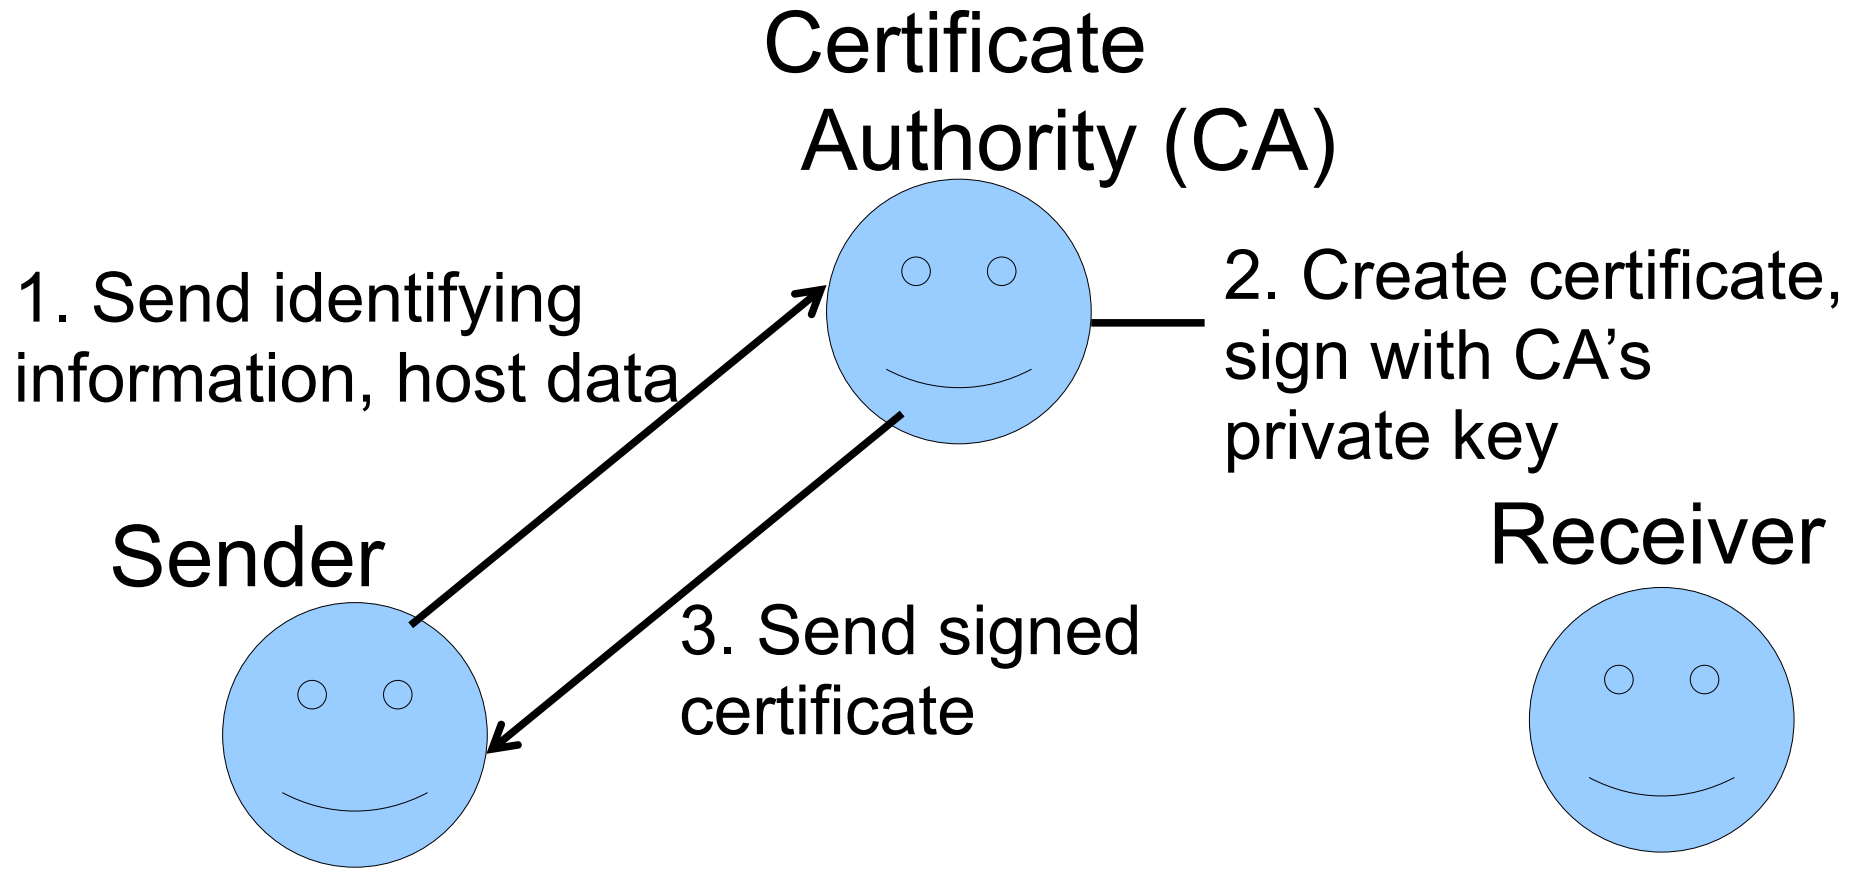


Certification authorities

- ▶ A certification authority is an organisation responsible for issuing and verifying the correctness of certificates
- ▶ If a host's certificate is signed by a CA, then any other host trusting the CA may reliably know that the certificate's public key belongs to the host as stated
- ▶ Publicly trusted CAs exist, such as VeriSign and CertCA




Applying for a certificate




Applying for a certificate

Certificate
Authority




Sender



1. Send digital certificate



Receiver



2. Verify signature
on certificate
with CA's public key



X.509

- ▶ X.509 is a popular form of certificate
- ▶ An X.509 certificate consists of three parts:
 - ▶ The certificate details
 - ▶ The signature of the certificate
 - ▶ The algorithm used to sign the certificate
- ▶ The certificate details then include:
 - ▶ A unique serial number for the certificate
 - ▶ The period (from X to Y) that the certificate is valid
 - ▶ The name of the certificate's issuer
 - ▶ A unique identifier for the issuer
 - ▶ The name of the certificate's owner
 - ▶ The public key of the owner



What we've covered

- ▶ Internet security threats
- ▶ HTTP authentication
- ▶ Digests and encryption
- ▶ Public key encryption
- ▶ Transport Layer Security and HTTPS
- ▶ Digital certificates

