

# Operating Systems Week 8

Linux

# Virtual machines

- ♦ Run a **guest os** on a **host os**

# VM Security

- ♦ Better than nothing

# Processes

- ♦ Each has, *inter alia*:
  - An owner (from last lecture: RUID, EUID, ...)
  - A parent process
  - A scheduler priority (‘niceness’)
- ♦ What happens when a process starts another?

# A shell (e.g. bash)

- ♦ What people think of as ‘a terminal’, or ‘the command line’
- ♦ Important notion: the **present working directory**, changed using **cd**.
- ♦ Most commands are executables on disk
  - e.g. /bin/ls; /bin/grep; /bin/more ...
  - A command starts a process, with some **arguments** (`String[] args`)
  - Process **returns an exit code**: 0 for success, non-zero otherwise. In Java - `System.exit(...)`.

# UNIX Philosophy

Courtesy of Mike Gancarz:

1. Small is beautiful.
2. Make each program do one thing well.
3. ...

# ConCATenate

- ♦ Loop over a list of files
  - Open;
  - Print out contents;
  - Close

# grep (**re** = regular expression)

- ♦ Take a regular expression
- ♦ Read some input
- ♦ Then:
  - By default, only print out lines that **match that regular expression**
  - `grep -v`: print out lines that do not match it
  - `grep -l`: print out filenames with a match
  - `grep -c`: count matching lines. `grep -cv?`



# **wc** – word count

- ♦ Loop over a list of files
  - Open
  - Count how many words/characters/lines
  - Close
  - Print out data
- ♦ Print total line across all files

# head, tail

- `head` – get the first lines:
  - Read 10 lines, print them out, stop
  - `head -lines=20` : read 20 lines then stop
- `tail` – get the last lines:
  - Read all the lines; once input has ended, print the last 10 out.
  - `tail -lines=20 ...`

# Running commands in sequence

- ♦ Semicolons ~= newlines

```
cat a.txt; wc b.txt; head c.txt
```


- ♦ Does not check the exit code was zero is at each stage – can be problematic e.g:

```
cd some_directory; rm a.txt
```

- ♦ Instead use &&:

```
cd some_directory && rm a.txt
```

# Processes



# System calls

Just like reading/writing to/from files:

```
write(1, "Hello World\n", 12);
```

```
write(2, "Uh-oh", 5);
```

```
char buf[256];
```

```
int bytesRead = read(0, buf, 256);
```

# I can haz Java?

Standard input = `System.in`

Standard output = `System.out`

Standard error = `System.err`

# Reading from System.in

- ♦ System.in is an **InputStream**

```
byte[] buf = new byte[256];  
int bytesRead =  
    System.in.read(buf, 0, 256);
```



Leads to a system call

# Unicode

- ♦ In Java: `chars` are 2 bytes each
  - There are more than 256 written characters in use in the world (no şit...)
- ♦ Unicode (in UTF-8 encoding) in a nutshell:
  - Read a byte. Easy case: the top-most bit is 0, turn into a char (following ASCII).
  - If the ***n*** top-most bits are 1, read another ***n* - 1** bytes, for ***n*** bytes total; combine into a char.
  - ***n*** can be up to 4.



# Reading chars from System.in

- ♦ This is exactly what an `InputStreamReader` is for:

```
InputStreamReader charsIn = new  
    InputStreamReader(System.in) ;  
char[] buf = new char[256] ;  
int charsRead = charsIn.read(buf,  
    0, 256) ;
```

# Back to system calls

- ♦ What if we write:

```
for (int i = 0; i < 256; ++i) {  
    int x = charsIn.read();  
    ...do something with x...  
}
```

# BufferedReader

```
BufferedReader b = new  
    BufferedReader(new  
        InputStreamReader(System.in) );
```

```
for (int i = 0; i < 256; ++i) {  
    int x = b.read();  
    ...do something with x...  
}
```

# Java OutputStream

```
OutputStream o = ...;
```

```
o.write(someBytes, 0, 256);
```

# Java BufferedOutputStream

Same as an OutputStream but with buffering:


```
OutputStream o = ...;
BufferedOutputStream b = new
    BufferedOutputStream(o);
for (int i = 0; i < 256; ++i) {
    b.write(...some byte...);
}
```

# System.out


- ♦ ...is a **PrintStream** – a **BufferedOutputStream** that also does char -> byte conversion

```
for (int i = 0; i < 256; ++i) {  
    System.out.print(...some char...);  
}
```

# Processes in a Terminal



# Pipes





# Example pipes

- ♦ How many lines in a file contain the word dave?

```
cat file.txt | grep 'dave' | wc -l
```

- ♦ Files whose names match the regexp 'steve\*', displaying results one page at a time:

```
ls | grep 'steve*' | more
```

# Inside the kernel

- ♦ Pipes = **producer-consumer**
  - **ls** is **producing** bytes into a buffer, every time it calls `write(1, ...)` ;
  - **grep** is **consuming** bytes from a buffer, every time it calls `read(0, ...)` ;
- ♦ Overkill solution: make the entire interrupt handler a **critical section**
  - In practice?

# /proc/....

- ♦ /proc is a **virtual filesystem** (see also /dev)
- ♦ /proc/12345 contains information about the process 12345. **Who owns /proc/12345?**
- ♦ /proc/12345/environ: environment variables
- ♦ /proc/12345/fd lists all the **file descriptors**
  - Using `ls -l /proc/12345/fd` can see what 0,1,2 are connected to

# /dev: Everything is a file

- ♦ /dev contains **hardware devices**
  - /dev/sda, /dev/sdb, /dev/sdc... - disks
  - /dev/snd/... - audio devices
  - /dev/video0 – a web cam
  - ...
- ♦ Opened/read from/written to, like files
- ♦ **Access restrictions:** ordinary users cannot read/write disks directly; groups for some devices (e.g. 'video' for video devices); ...

# Redirection

- ♦ Produce a file containing all the lines from three input files, that match 'A[a-z]\*A':

```
cat 1.txt 2.txt 3.txt |  
grep 'A[a-z]*A' > matched.txt
```

**NB Do not confuse > and |**

# Redirection to /dev/...

Write a disk image:

```
cat disk.img > /dev/somedisk
```

Run a command, throw away output:

```
command > /dev/null
```

Send random input to a command:

```
cat /dev/random | somecommand
```

# Buffering and redirection

```
some-command > output.log
```

- ♦ Q: How often is output.log updated?
- ♦ A: When the buffer fills, or flush() is called.
- ♦ Can we turn-off the buffer?

```
stdbuf -o 0 <command> : runs <command> with  
0-byte output buffer
```

```
stdbuf -oL <command> : flush at newlines
```

# stdout, stderr

- ♦ By default, stdout is redirected, not stderr
- ♦ Can use 2 to redirect stderr:

```
cat a.txt b.txt > joined.txt 2> error.log
```

```
cat a.txt b.txt 2> error.log | wc -l
```

- ♦ Can also merge stderr with stdout:

```
cat a.txt 2>&1
```



# More commands: sed

- ♦ sed: **s**tream **e**ditor
- `sed -e 's,cat,dog,'`
  - Replace cat with dog, once
- `sed -e 's,cat,dog,g'`
  - Replace cat with dog, **globally**
- `sed -e 's,\([a-z]+\)[0-9]*,\1,'`

# find

- ♦ Finds files with certain properties, recursively
- `find -name "foo"`: find files named foo
- `find -iname "*.jpg"`: find jpg files, case insensitive (will match .JPG, .jpg...)
- `find -type d`: find all directories (`-type f` for all files)


# find ... -exec

- ♦ -exec: run a command on each file found

```
find -iname "*.txt" -exec cat {} \;
```



placeholder for  
the filename



end of the  
command  
to run on  
the file

# Putting it together

- ♦ A directory tree contains CSV files of exam results – one CSV per module, giving the mark for each candidate.
- ♦ The anonymous marking code T01234 is for the student Geoff Vader.
- ♦ Make a file results.csv, containing all results for T01234; and replacing the string “T01234” with the string “Geoff Vader”

# ps: process status

- ♦ ps: shows processes running in the current terminal
- ♦ ps ux: show all the current **user**'s processes, in **e**x~~x~~ended detail
- ♦ ps aux: same but for **a**ll users
- ♦ ps auxf: ASCII-art process tree
- ♦ What can you see on the right on ps aux?

# Security?

somecommand

--username=derp

--password=my\_super\_secret\_password

# Process signals: kill

- `kill 28125`: terminate process 28125
- ♦ Sends `SIGTERM` – can optionally be caught, and terminate ‘nicely’ before closing
- `kill -SIGKILL 28125`: sends `SIGKILL`, killing it immediately without question
- `kill -SIGSTOP 28125`: pause 28125. (It is never chosen by the scheduler.)
- `kill -SIGCONT 28125`: resume 28125

# Hangup

- ♦ The **hangup** signal is sent when the **terminal is disconnected** (e.g. window is closed)
- ♦ Or, `kill -SIGHUP 28125`
- ♦ Usually the same as `kill -SIGTERM`
- ♦ Or can ignore it using `nohup`:  
`nohup somecommand`



# Returning to the shell

- `cat a.txt b.txt c.txt`: cats three files, then returns to the shell prompt ready for the next command.
- `cat a.txt b.txt c.txt &` will return to the shell immediately.

`nohup web-server >& output.log &`

# Priority

- ♦ Each process has a niceness. Default = 0
  - 20 = only gets 'idle' CPU cycles
  - -20 = highest possible priority

`/usr/bin/nice <command>`: runs `<command>`  
with nice 10

`/usr/bin/nice -n 20 <command>`: runs  
`<command>` with nice 20

`renice -n 20 28125`: set PID 28125 to nice 20

# The rules of nice

- ♦ Processes start with nice 0
- ♦ **root** can give processes any nice value, from -20 to 20
- ♦ **You** can decrease **your** processes' priorities: bigger nice values
- ♦ You **cannot increase your processes' priorities** as that might undermine root's authority

# Environment variables

- ♦ A map from strings to values
- PWD: the **present working directory**  
**echo \$PWD**
- PATH: where to look for commands
  - e.g. /usr/sbin:/usr/bin:/sbin:/bin
- HOME: path to the home directory
- LANGUAGE: current language
- ♦ ...

# cd

- ♦ A shell **built-in**: changes the present working directory. Also:
  - `pushd <dir>`: pushes PWD onto a stack; then does `cd <dir>`
  - `popd`: pop PWD off a stack
- ♦ Why is `cd` built into the shell? Could `cd` be an executable, e.g. saved as `/bin/cd`?

# Scripting

- ♦ Writing magic sequences of commands, piped together, is tedious – write a script
- ♦ Simple case: just write the commands
- ♦ The shell also has loops, variables, ...
- ♦ Could use a scripting language, e.g. perl
- ♦ If possible, write your scripts as **filters**: read from stdin, write to stdout